



Εθνικό Μετσόβιο Πολυτεχνείο

Σχολή Ηλεκτρολόγων Μηχανικών
και Μηχανικών Υπολογιστών

Τομέας Τεχνολογίας Πληροφορικής
και Υπολογιστών

Implementation of an $O(1)$ in-memory cache for a distributed storage service

ΔΙΠΛΩΜΑΤΙΚΗ ΕΡΓΑΣΙΑ

ΟΝΟΜΑ ΦΟΙΤΗΤΗ

Επιβλέπων : Υπεύθυνος Διπλωματικής

Τίτλος Υπευθύνου

Αθήνα, Σεπτέμβριος 9999



Εθνικό Μετσόβιο Πολυτεχνείο

Σχολή Ηλεκτρολόγων Μηχανικών
και Μηχανικών Υπολογιστών

Τομέας Τεχνολογίας Πληροφορικής
και Υπολογιστών

Implementation of an $O(1)$ in-memory cache for a distributed storage service

ΔΙΠΛΩΜΑΤΙΚΗ ΕΡΓΑΣΙΑ

ΟΝΟΜΑ ΦΟΙΤΗΤΗ

Επιβλέπων : Υπεύθυνος Διπλωματικής

Τίτλος Υπευθύνου

Εγκρίθηκε από την τριμελή εξεταστική επιτροπή την 9η Σεπτεμβρίου 9999.

.....
Πρώτο μέλος επιτροπής
Τίτλος μέλους

.....
Δεύτερο μέλος επιτροπής
Τίτλος μέλους

.....
Τρίτο μέλος επιτροπής
Τίτλος μέλους

Αθήνα, Σεπτέμβριος 9999

.....
Όνομα Φοιτητή

Διπλωματούχος Ηλεκτρολόγος Μηχανικός και Μηχανικός Υπολογιστών Ε.Μ.Π.

Copyright © Όνομα Φοιτητή, 9999.

Με επιφύλαξη παντός δικαιώματος. All rights reserved.

Απαγορεύεται η αντιγραφή, αποθήκευση και διανομή της παρούσας εργασίας, εξ ολοκλήρου ή τμήματος αυτής, για εμπορικό σκοπό. Επιτρέπεται η ανατύπωση, αποθήκευση και διανομή για σκοπό μη κερδοσκοπικό, εκπαιδευτικής ή ερευνητικής φύσης, υπό την προϋπόθεση να αναφέρεται η πηγή προέλευσης και να διατηρείται το παρόν μήνυμα. Ερωτήματα που αφορούν τη χρήση της εργασίας για κερδοσκοπικό σκοπό πρέπει να απευθύνονται προς τον συγγραφέα.

Οι απόψεις και τα συμπεράσματα που περιέχονται σε αυτό το έγγραφο εκφράζουν τον συγγραφέα και δεν πρέπει να ερμηνευθεί ότι αντιπροσωπεύουν τις επίσημες θέσεις του Εθνικού Μετσόβιου Πολυτεχνείου.

Περίληψη

TODO: Φτιάξε την ελληνική περίληψη και τις λέξεις κλειδιά

Λέξεις κλειδιά

Λέξη-κλειδί 1, λέξη-κλειδί 2, λέξη-κλειδί 3

Abstract

TODO: Write the english summary and the keywords

Key words

Key-word 1, Key-word 2, Key-word 3

Ευχαριστίες

TODO: Thank everyone

Όνομα Φοιτητή,
Αθήνα, 9η Σεπτεμβρίου 9999

Η εργασία αυτή είναι επίσης διαθέσιμη ως Τεχνική Αναφορά CSD-SW-TR-*-* , Εθνικό Μετσόβιο Πολυτεχνείο, Σχολή Ηλεκτρολόγων Μηχανικών και Μηχανικών Υπολογιστών, Τομέας Τεχνολογίας Πληροφορικής και Υπολογιστών, Εργαστήριο Τεχνολογίας Λογισμικού, Σεπτέμβριος 9999.

URL: <http://www.softlab.ntua.gr/techrep/>

FTP: <ftp://ftp.softlab.ntua.gr/pub/techrep/>

Contents

| | |
|--|----|
| Περίληψη | 5 |
| Abstract | 7 |
| Ευχαριστίες | 9 |
| Contents | 11 |
| List of Figures | 15 |
| List of Tables | 17 |
| 1. Introduction | 21 |
| 1.1 Thesis subject | 22 |
| 1.2 Thesis structure | 22 |
| 2. Necessary theoretical background | 23 |
| 2.1 Multithreading | 23 |
| 2.2 Interprocess Communication - IPC | 24 |
| 2.2.1 Signals | 24 |
| 2.2.2 Sockets | 25 |
| 2.2.3 Shared memory | 27 |
| 2.3 Concurrency control | 27 |
| 3. Archipelago | 29 |
| 3.1 Overview | 29 |
| 3.2 Architecture | 29 |
| 4. Scalability, Tiering and Caching | 33 |
| 4.1 What is scalability? | 33 |
| 4.2 What is tiering? | 33 |
| 4.3 What is caching? | 34 |
| 4.3.1 Write policies | 35 |
| 4.3.2 Caching limitations | 35 |
| 4.4 Real-life scenario | 36 |
| 4.5 Current Solutions | 37 |
| 4.5.1 Beache | 37 |
| 4.5.2 Flashcache | 38 |
| 4.5.3 EnhanceIO | 38 |
| 4.5.4 Memcached | 39 |
| 4.5.5 Couchbase Server | 40 |
| 4.5.6 Honorary mentions | 40 |
| 4.5.7 Evaluation | 41 |

| | |
|--|----|
| 5. Design of cached | 43 |
| 5.1 Design rationale | 43 |
| 5.2 Cached components | 44 |
| 5.2.1 Overview | 45 |
| 5.2.2 The xcache xtype | 45 |
| 5.2.3 xcache flow | 48 |
| 5.2.4 The xworkq xtype | 48 |
| 5.2.5 The xwaitq xtype | 49 |
| 5.3 Cached Design | 49 |
| 5.3.1 Request handling | 50 |
| 5.3.2 Write policy enforcing | 51 |
| 5.3.3 Data propagation | 51 |
| 5.3.4 Bucket pool | 51 |
| 5.3.5 Request wrapper (Cache-IO) | 52 |
| 5.3.6 Book-keeping utilities | 52 |
| 5.4 Cached Flow | 53 |
| 5.4.1 Write | 53 |
| 5.4.2 Read | 53 |
| 5.4.3 Copy | 53 |
| 5.4.4 Info | 53 |
| 6. Implementation of cached | 55 |
| 6.1 Implementation of xcache | 55 |
| 6.1.1 xcache initialization | 55 |
| 6.1.2 Cache entry preallocation | 56 |
| 6.1.3 Cache entry initialization | 57 |
| 6.1.4 Cache entry indexing | 58 |
| 6.1.5 Entry eviction | 59 |
| 6.1.6 Entry reinsertion | 60 |
| 6.1.7 Concurrency control | 60 |
| 6.1.8 Event hooks | 62 |
| 6.2 Implementation of cached | 62 |
| 6.2.1 Cached initialization | 63 |
| 6.2.2 Bucket pool | 65 |
| 6.2.3 Request handling | 66 |
| 6.2.4 "ENOSPC" scenarios | 66 |
| 7. Performance evaluation of cached | 67 |
| 7.1 Benchmark methodology | 67 |
| 7.2 Specifications of test-bed | 68 |
| 7.3 Performance comparison between cached and sosd | 68 |
| 7.3.1 Workload smaller than cache size - Peak behavior | 69 |
| 7.3.2 Workload larger than cache size - Sustained behavior | 73 |
| 7.4 Performance evaluation of cached parameters | 75 |
| 8. The synapsed peer | 79 |
| 8.1 Design of synapsed | 79 |
| 8.2 Implementation of synapsed | 80 |
| 8.2.1 Synapsed initialization | 80 |
| 8.2.2 Request polling | 80 |
| 8.2.3 Marshalling | 81 |
| 8.2.4 Send/Receive | 82 |

| | |
|--------------------------------------|----|
| 8.3 Evaluation of synapsed | 82 |
| Bibliography | 83 |

List of Figures

| | | |
|------|---|----|
| 2.1 | Data encapsulation for a UDP packet | 26 |
| 3.1 | Archipelago overview | 29 |
| 3.2 | Archipelago components | 30 |
| 4.1 | Computer Memory Hierarchy | 34 |
| 5.1 | Cached design | 50 |
| 7.1 | Comparison of bandwidth performance for writes | 69 |
| 7.2 | Comparison of bandwidth performance for reads | 70 |
| 7.3 | Comparison of latency performance for writes | 72 |
| 7.4 | Comparison of latency performance for reads | 72 |
| 7.5 | Comparison of bandwidth performance for writes | 73 |
| 7.6 | Comparison of bandwidth performance for reads | 74 |
| 7.7 | Comparison of latency performance for writes | 74 |
| 7.8 | Comparison of latency performance for reads | 75 |
| 7.9 | Bandwidth performance per number of threads | 76 |
| 7.10 | Latency performance per number of threads | 76 |
| 7.11 | Latency performance of cold/warm cache for variable sizes | 77 |

List of Tables

| | | |
|-----|---|----|
| 4.1 | Access times of storage mediums | 34 |
| 6.1 | Command line arguments of cached | 63 |
| 7.1 | Test-bed hardware specs | 68 |
| 7.2 | Test-bed software specs | 68 |
| 7.3 | Write performance results for 2-threaded cached | 71 |

List of Listings

| | | |
|------|--|----|
| 6.1 | <code>xcache_init</code> definition | 55 |
| 6.2 | Main <code>xcache</code> struct | 56 |
| 6.3 | <code>xcache</code> struct fields for preallocated entries | 56 |
| 6.4 | <code>xcache_entry</code> struct | 56 |
| 6.5 | <code>xcache_entry</code> fields, relevant for preallocation | 57 |
| 6.6 | Cache entry allocation/initialization function | 57 |
| 6.7 | <code>xcache</code> struct fields for entry indexing | 58 |
| 6.8 | Indexing functions | 58 |
| 6.9 | <code>xcache_entry</code> struct relevant indexing | 58 |
| 6.10 | <code>xcache</code> struct fields for eviction | 59 |
| 6.11 | Doubly-linked LRU list | 59 |
| 6.12 | Concurrency control fields | 60 |
| 6.13 | Atomic operations of GCC | 61 |
| 6.14 | <code>xcache_ops</code> struct | 62 |
| 6.15 | Main cached struct | 63 |
| 6.16 | Cached entry struct | 64 |
| 6.17 | Bucket implementation | 65 |
| 8.1 | Synapsed header | 81 |
| 8.2 | GCC <code>pragma pack</code> directive | 82 |

Chapter 1

Introduction

FIXME: Polish text below

In the new-age race of hardware, where vendors are always trying to produce the next record-breaking unit, be it CPU, RAM or network cards, hard disks seem old and gasping to catch up. For decades now, their sub-par performance, compared to volatile memory, has been the bottleneck of every IO-intensive application and the headache of storage designers [Caul10]. As over-stretched as this may seem, it is a fact that it has shaped the way storage is built; from the common computer (Linux's page cache, Bcache) to large deployments (the memcached servers of Facebook and Twitter **FIXME:** add paper), there is a tremendous effort from the research and corporate community that is being invested in sidestepping hard disks and finding alternative methods to store data.

The hard disks' industry answer to this is the continuous drop of their prices. In 2011 (**FIXME:** be more specific), the HDDs reached their all-time low price of \$0.053/GB [hdd-]. Moreover, the emerging movement of greener data centers has benefited hard disks, since their low energy costs has been proved attractive to the enterprises. Yet, for how long can the HDD industry keep lowering the costs to mitigate their lack of performance?

The answer came very fast and unfortunately in a tragic way. The end of July of 2011 marked the beginning of a 6-month turmoil for Thailand, with a flood that was described as "the worst flooding yet in terms of the amount of water and people affected" [floo]. The hard disk industry also suffered a huge hit due to the fact that 25% percent of the global hard disk production was from factories in Thailand, that were also affected by the flood.

The result was an overnight 40% percent increase of hard disk prices. The reasons behind this increase were partly to compensate for the flood damages and partly to seize the opportunity to increase the profit margins of the two biggest producers, Western Digital and Seagate, from 6% and 3% to 16% and 37% respectively [Rose12].

The timing could not have been worse for the HDD industry. This increase in price introduced SSDs who are now starting to be considered as a viable solution for peripheral storage tasks such as journaling and caching, due to their high performance and their persistence, which separates them from other volatile storage types.

On the other hand, the current situation is that HDDs can only marginally improve their performance. As their rotational speed approaches the speed of sound, their production will be rendered at best difficult, and their heat generation, power consumption and lack of long-term reliability will make their adoption prohibitive [Seic06],[soun].

The cloud world is largely affected by the future of storage mediums. **FIXME:** explain why. Besides SSDs, there are various flash memory types such as the IOdrive of Fusion IO that are being utilized in performance-intensive environments, albeit for higher prices. While this tug-of-war continues between flash memory and HDDs and there is no clear winner, clever designs and buffering techniques are very important for software services.

The Synnefo [?] cloud software, which powers the okeanos cloud service [?] **FIXME:** explain how it is affected and where Archipelago enters.

1.1 Thesis subject

TODO: explain that we are going to create a cache which will reside in memory and see if it can help with this situation

1.2 Thesis structure

FIXME: Fix the summary and references of these chapters.

Chapter ??: We define what "cloud" means and mention some of the most notable examples. Then, we give a brief overview of the synnefo implementation, its key characteristics and why it can have a place in the current cloud world.

Chapter 3: We present the architecture of Archipelago and provide the necessary theoretical background (mmap, IPC) the reader needs to understand its basic concepts. Then, we thoroughly explain how Archipelago handles I/O requests. Finally, we mention what are the current storage mechanisms for Archipelago and evaluate their performance.

Chapter ??: We explain why tiering is important and what is the state of tiered storage at the moment (bcache, flashcache, memcached, ramcloud, couchbase). Then, we provide the related theoretical background for cached (hash-tables, LRUs). Finally, we defend why we chose to roll out our own implementation.

Chapter 5: We explain the design of cached, the building blocks that is consisted of (xcache, xworkq, xwaitq). Then, we give some examples that illustrate the operation under different scenarios

Chapter 6: We present the cached implementation, the structures that have been created and the functions that have been used.

Chapter 7: We explain how cached was evaluated and present benchmark results.

Chapter ??: It connects brain parts. And its tale must be told.

Chapter ??: We draw some concluding remarks and propose some future work.

Chapter 2

Necessary theoretical background

FIXME: Make this section about the concepts that will be discussed, as well as point to the appropriate chapters.

2.1 Multithreading

Multithreading is a programming concept that has been the subject of research long before the emergence of SMP systems ¹. More specifically, temporal multithreading has been introduced in the 1950s whereas Simultaneous Multithreading (SMT), which is the current invocation of multithreading programming, was first researched by IBM in 1968[mt].i **TODO:** Add what is temporal mt and simultaneous mt

Before threads, programs could utilize the concurrency of SMP systems using forked processes that would communicate with each other. The introduction of threads did not render this practice obsolete, but instead provided an alternative technique to speed up applications.

Threads and process have some fundamental differences, which are shown in the following list:

- Threads are always parts of a process, whereas processes are independent from each other and may only have a parent-child connection between them.
- Forked processes have their own address space and resources, which are inherited by the parent process with CoW semantics. Multiple threads, on the other hand, usually share the same memory and resources with the other threads in the same process.

From the above differences, we can see that there are no clear advantages of a multithreaded approach over a multiprocess one. To better demonstrate our point, we will present the advantages and disadvantages of multithreading programming in the following lists:

The advantages are:

- Context switching is generally faster between threads, mostly due to the fact that the TLB ² cache does not need to be flushed. The TLB cache misses are expensive and are avoided as much as possible[Barr10].
- Sharing data between threads is easier, due to the fact that they use the same memory by default.

¹ Symmetric multiprocessing systems, commonly systems with multiple processors

² Translation Lookaside Buffer, a hardware cache that speeds up the translation of virtual addresses to physical RAM pages.

Whereas the disadvantages are:

- Processes are more isolated than threads, which means they are guarded against two things: i) thread-unsafe functions and ii) data corruptions, which if they happen to one thread they bring the whole process to a halt.

Regardless of the chosen method, at some point the programmer will have face two of the biggest challenges of multithreading/multiprocess programming; interprocess communication, which is discussed in Section 2.2, and concurrency control, which is discussed in Section 2.3.

2.2 Interprocess Communication - IPC

Interprocess Communication is a concept that predates the SMP systems that we all use nowadays. It is a set of methods that an OS uses to allow processes and threads to communicate with each other. Archipelago for example, uses various IPC methods to synchronize its different components.

The full list of Linux's IPC methods is presented below:

- *Signals*, which are sent to a process to notify it that an event has occurred.
- *Pipes*, which are a one-way channel that transfers information from one process to another.
- *Sockets*, which are bidirectional channels that can transfer information between two or more processes either locally or remotely through the network.
- *Message queues*, which is an asynchronous communication protocol that is used to exchange data packets between processes.
- *Semaphores*, which are abstract data types that are used mainly for controlling accesses on a same resource.
- *Shared memory*, which is a memory space that can be accessed and edited by more than one process.

We will concentrate on the following IPC methods: i) signals, ii) sockets and iii) shared memory, since these are the methods that Archipelago and our implementation use.

2.2.1 Signals

Signals are notifications that are sent to processes and can be considered as software interrupts. The signal's purpose is to interrupt the execution of a process and inform it that an event has occurred.

Given that there more than one events and exceptions that can occur in a system, there are also various signals that match to each one of these events. For more information about the signals that Linux supports as well as the conditions on which they are raised, the reader is prompted to consult the man pages for `signal(7)` or read the POSIX.1-1990, SUSv2 and POSIX.1-2001 standards.

Moreover, the above standards dictate the standard behavior of a process when a signal is received. The standard actions that a process can take, fall roughly in the following categories:

- ignore the signal,

- pause its execution,
- resume its execution or
- stop its execution and/or dump its core

Finally, a process is not limited to this set of actions. It can instead do one of the following things for each signal, with the exception of SIGKILL and SIGSTOP signals:

- ignore the signal
- block the signal, which is part of the Archipelago IPC and its usage is described in Section ?
- install a custom signal handler function, which essentially passes the signal handling task to the process.

2.2.2 Sockets

Sockets are a bidirectional means **FIXME:** (is it means or mean ?) of sending data between processes. The processes can be in the same host but most commonly, they are in remote hosts and the data are sent over the network. Furthermore, from all the IPC methods that we have described above, sockets are the only method that enables remote communication.

There are many socket implementations for different purposes, which are divided in several communication domains, most of which are rather obscure. The three communication domains, however, that are supported by most UNIX and UNIX-like operating systems are:

- *IPv4* domain, which allows communication between processes over the Internet Protocol version 4 network.
- *IPv6* domain, which allows communication between processes over the Internet Protocol version 6 network.
- *UNIX* domain, which allows communication between processes in the same host

The above three communication domains are further divided in two types, based on the transport layer protocol that they use.

- *Stream sockets*, which use the Transmission Control Protocol (TCP) or Stream Control Transmission Protocol (SCTP),
- *Datagram sockets*, which use the User Datagram Protocol (UDP),

The TCP/UDP protocols are only one layer out of the four layers of the TCP/IP protocol stack that the RFC 1122[1122] defines, and we will explain them in detail in the following sections. Although a thorough explanation of the TCP/IP protocol stack is out of the scope of this thesis and is not needed to understand the following sections, we will provide a brief explanation of it for the sake of completeness.

The TCP/IP protocol stack is the basis for the World Wide Web and the most used form of networking. It specifies all the stages of the data processing that need to happen in various levels and entities, such

as operating systems, network cards, routers etc. in order to connect two machines over the network. For this reason, the data that are sent are encapsulated in layers, which can be seen in Figure 2.1.

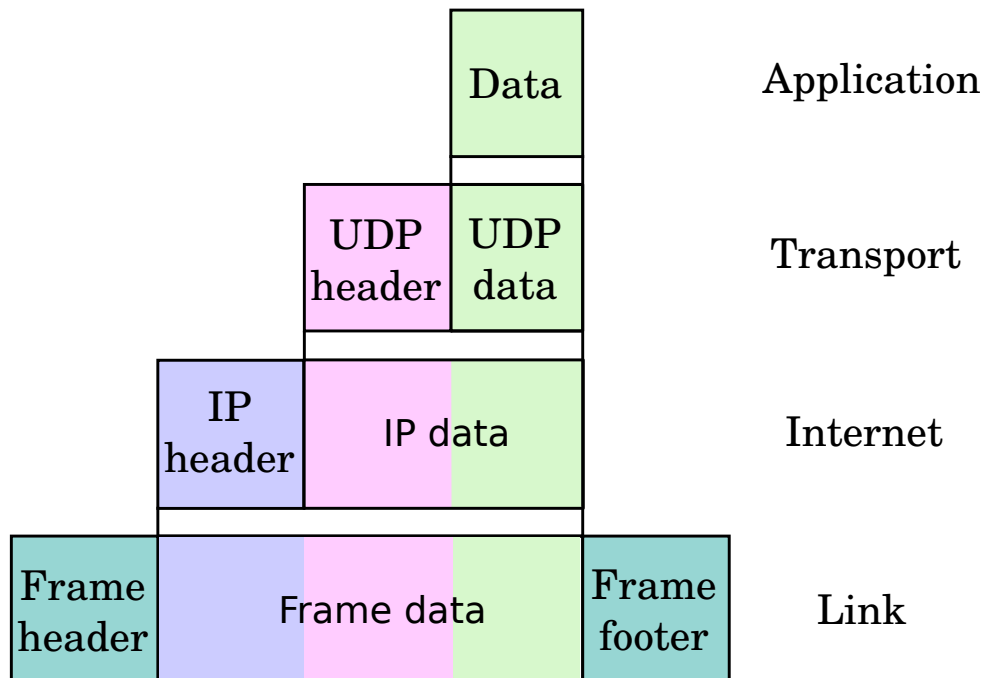


Figure 2.1: Data encapsulation for a UDP packet

We now continue with an presentation of TCP and UDP protocols.

TCP

The Transmission Control Protocol is connection-oriented, i.e. it provides unique connection between two sockets, and has the following key features:

Reliability The data will arrive to the receiver as a whole, or they will not arrive at all. In the latter case, the receiver may receive spurious packets but it will not acknowledge them until it has received all of them.

Ordered transfer The data will arrive in the same order that they were sent.

Error-checking The data are checksummed to allow the receiving end to check if there was any data corruption.

Rate-limiting When the receiver accepts packets with slower rate than the sender, the sender will adjust its rate to ensure packet delivery and less congestion.

Byte-stream The data that are sent do not have a boundary.

UDP

FIXME: Make this a list The User Datagram Protocol on the other hand has far less restrictions than TCP. First, it is connectionless, meaning that the socket can receive requests from anyone. Second,

it provides no guarantees about the delivery of the messages. Third, the messages can arrive in other order than the one they were sent. Fourth, there is no rate-limiting, meaning that the congestion control must be handled in the application level. Finally, it sends datagrams instead of bytes, which have definite boundaries.

Although there might seem that there is no reason for one to choose the unreliable UDP protocol, the lack of the TCP overhead makes it an ideal choice for applications that value speed over packet loss.

2.2.3 Shared memory

When two or more processes share the same memory segment, they can exchange data by placing it in a region of the segment. The data then becomes instantly visible to the other processes too, since their page-table entries for this segment point to the same physical RAM pages.

A popular way of mapping shared memory to a process's address space, which is also used in Archipelago, is with POSIX `mmap()`. There are two mapping types of mapping:

- *Private mapping*, in which case the mapping contents will not be visible to other processes that have mapped the same file and
- *Shared mapping*, in which case the mapping contents will be visible to all processes that map this file and changes to the mapping will be propagated to the shared memory.

Finally, an issue with mappings is that the start of the shared memory is not always mapped in the same virtual address for all processes. For this reason, when processes want to share data, they should not pass direct pointers to them, but relative pointers (i.e. offsets) from the start of the segment, which are common for all processes and can be translated to the correct direct pointers.

2.3 Concurrency control

Concurrency control is the set of methods that a program uses to ensure that concurrent accesses to the same data will leave them in a consistent state.

There are several techniques that are used for concurrency control and are listed below:

- *Spinlocks*, which are locks that protect a critical segment. Typically, a thread acquires a lock at the start of the critical segment and releases it at the end of it. Threads that are waiting for the lock essentially "spin", i.e. they busy-loop until the lock is released.
- *Mutexes*, which are locks that protect a critical segment in the same fashion as spinlocks. Their difference from spinlocks, however is that if a thread cannot get the lock, it will block instead of busy-loop.
- *Semaphores*, which are also an IPC method. In concurrency control context, they are abstract data types that restrict the number of simultaneous accesses to a resource or a critical segment. When the number of times is one (1), they essentially degenerate to mutexes, with the main difference that they have no concept of an owner.
- *Atomic operations*, which are hardware-assisted operations whose purpose is to atomically update a value as fast as possible. The atomicity is usually achieved with implicit hardware locks on the bus or cache-line. Atomic operations come in many flavors such as "add-and-fetch", "compare-and-swap" etc.

Concurrency control - and locking in particular - have three important caveats that the programmer needs to know before he/she decides on the techniques that will be used:

Lock overhead

Lock overhead is the overhead that the locking mechanism introduces. For example, semaphores are a mechanism with big overhead, since they must be read and written to using system calls. If the critical segment they protect is simply the update of a variable, then the programmer is probably better off using a spinlock or atomic operations.

Lock contention

Lock contention can be considered as the overhead of the coarseness of the lock. There is contention for a lock when it is requested by many threads, to the point that the waiting time is longer than the execution time. This has a big performance impact to the implementation, since threads may consume their scheduled time spinning until they acquire a lock, or sleeping while they could do something more useful. The solution to this problem is commonly to redesign the locking scheme in order to break such locks into smaller ones.

Deadlocks

Deadlock is a situation in a multi-lock scenario where each process in a group of processes needs to acquire a lock which is held by another process in the same group. Since no process can continue, the operation of the group is essentially stalled. As a rule of thumb, the circular dependency of the locks can break if the locks are acquired in a predefined order. This however is only possible in less complex scenarios and in general, a more well-thought design is required.

TODO: Add networking and explain how we write to/from sockets, how we connect to sockets, what is polling, what is marshaling

TODO: Add section that explains the various indexing mechanisms, amongst which the hash tables and b-trees

Chapter 3

Archipelago

TODO: Add intro that explains what we will see as well as the Sections

3.1 Overview

Archipelago is the Volume Service of the Synnefo cloud software. It is responsible for creating Copy-on-Write, snapshottable volumes for VMs. Archipelago can be considered as a storage layer (see Figure 3.1) that is positioned between the VM's block device, and the underlying storage.




Figure 3.1: Archipelago overview

Archipelago has the following objectives:

- Thinly provision volumes to VMs with zero data movement.
- Snapshot VM volumes and use them as system images with, again, zero data movement.
- Allow VM migrations between Archipelago nodes with no restrictions.
- Be agnostic to the actual storage backend used.

3.2 Architecture

Archipelago has a modular architecture, which allows it to categorize its operations and assign them to distinct components. The IPC between these components, which are called **peers** in the Archipelago dialect, is facilitated by XSEG.

XSEG can be considered as the kernel of Archipelago and is a custom mechanism that: i) defines a common communication protocol for all peers, regardless of their type (userspace/kernelspace, singlethreaded/multithreaded) and ii) builds a shared memory segment, where peers can share data using zero-copy techniques. The above are provided to the peers by the `libxseg` library.

In Figure 3.2, we present the architecture of Archipelago. Moreover, we show the Archipelago peers, the communication channels between them and we briefly explain the operations that they are responsible for.

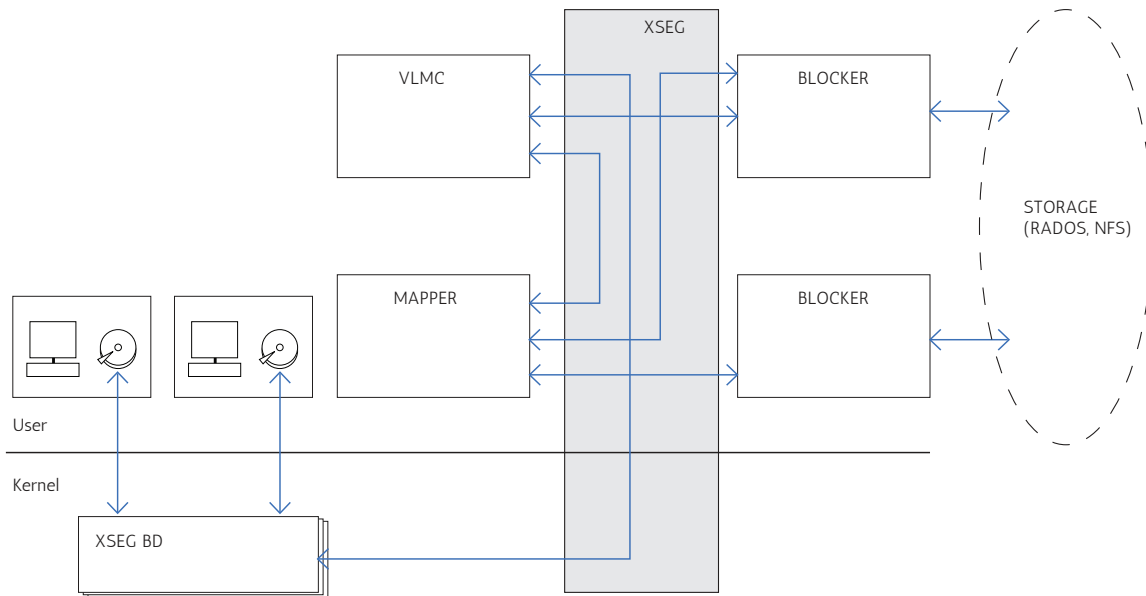


Figure 3.2: Archipelago components

XSEG Block Device (xsegbd)

`xsegbd` is a kernel module that exposes a VM's volume as a block device. This block device provides the entry point for the requests that enter the Archipelago layer.

VoLuMe Composer Daemon (vlmcd)

`vlmcd` accepts requests from the various `xsegbd` devices and translates them to object requests, with the help of `mapperd`.

Mapper Daemon (mapperd)

`mapperd` is responsible for the mapping of volumes to objects. This means that it must tackle a broad set of tasks such as knowing the objects that a volume consists of, cloning and snapshotting volumes and creating new ones.

File Blocker Daemon (blockerd)

`blockerd` is not a specific entity but a family of drivers, each of which is written for a specific storage type. File blockers have a single purpose, to read/write objects from/to the storage. Currently, there are file blockers for NFS and the RADOS object storage.

TODO: Add how a request is handled by archipelago (e.g. it goes to `xsegbd`, then `vlmcd`, then `mapperd`, etc.)

TODO: Add a generic description of how Archipelago handles the polling of requests (e.g. it binds ports, it waits for requests in them until it gets signalled. It uses polling for a certain number of cycles to solve bursts etc.

TODO: Add RADOS definition, cite philipgian and explain that sosd is used for...

Chapter 4

Scalability, Tiering and Caching

In this chapter, we will discuss the challenges of today's data storage and will attempt to explain the role of scalability, tiering and caching in mitigating costs and increasing performance. Moreover, we present the current solutions for boosting performance and we evaluate if they can be used in conjunction with Archipelago.

The structure of this chapter is the following. Sections 4.5, 4.2 and ?? explain what scalability, tiering and caching mean respectively. Section 4.4 attempts to exhibit the need for these techniques by providing a typical real-life scenario in which they can be used. Finally, Section ?? lists and evaluates some of the 3rd party, open-source solutions that employ the aforementioned techniques.

4.1 What is scalability?

Scalability, in storage service context, is the ability of the service to achieve two specific things:

1. accommodate the growth of load in a manner that does not impact the quality of the service and
2. utilize the addition of new resources to their full extend, in order to improve its performance.

There are two methods of scaling, horizontal (scaling out) and vertical (scaling up), which are explained below:

- *Horizontal scaling* applies to distributed services. It relies on the principle that adding more nodes to a system will mitigate the high load of the other nodes.
- *Vertical scaling* applies to all types of systems and refers to the addition of more resources such as better hardware, more RAM etc. to a node of the system.

The rule of thumb about these methods is that scaling up is the simpler solution for a service, albeit its performance cannot be increased much due to hardware limitations. On the other hand, scaling out is far more complex and requires a robust method of managing many nodes as well as their failures, but it may have lower costs (if nodes are made of commodity hardware), and has theoretically no limitations in performance gain (especially in share-nothing architectures).

4.2 What is tiering?

Tiering is the organization of different storage types in levels (or tiers) depending on their performance. These storage types usually differ in one of the following attributes: capacity, price or performance.

| Medium | Access time (ns) |
|-------------------------|------------------|
| CPU registers and cache | < 10 |
| RAM | $< 10^2$ |
| SSD | $< 10^5$ |
| Hard Disk | $< 10^7$ |

Table 4.1: Access times of storage mediums

Tiers such as SSD arrays or caches are necessary in most medium or larger deployments, in order to bridge the performance gap between RAM and magnetic disks, which can be seen in Table 4.1. To understand the need for tiering, consider the fact that when data do not reside in RAM and SSDs are not used, the performance penalty is $\times 10,000$ times the access time of RAM.

Tiered storage is analogous to the computer architecture model of memory hierarchy, which can be seen in Figure 4.1. Tiered storage is based on the same principles as memory hierarchy, in the sense that its objective is to keep "hot" data, i.e. data that are requested frequently, in the higher tiers.

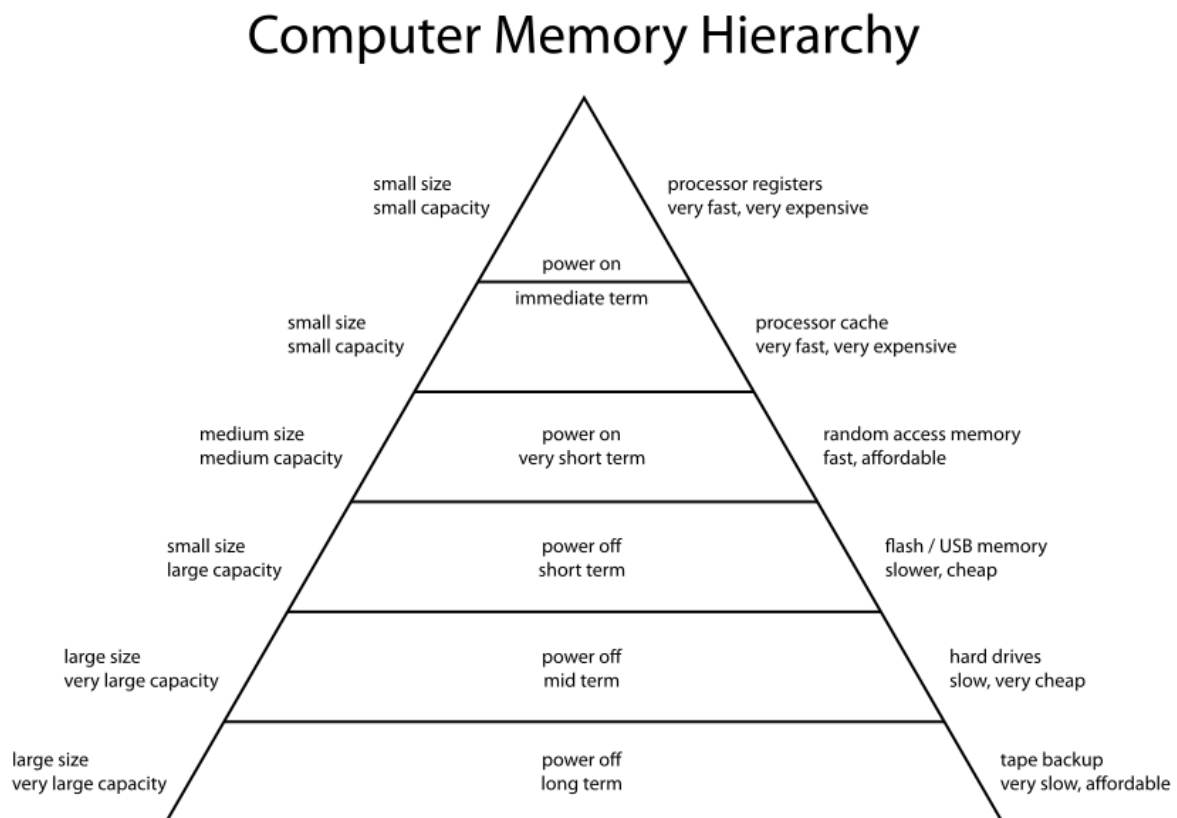


Figure 4.1: Computer Memory Hierarchy

4.3 What is caching?

In the context of I/O requests, caching is the addition of a fast medium in a data path, whose purpose is to transparently store the data that are intended for the slower medium of this data path. The benefits from caching is that later accesses to the same data will be faster than fetching them from the slower medium.

Caching is extensively used in computer architecture, as is evident from Figure 4.1. Besides the memory hierarchy model, it is a widely employed concept in storage services where fast mediums are used as journals for slower ones, essentially caching the data, or where dedicated servers are used to cache data, like in memcached's case (see more in Section 4.5.4).

4.3.1 Write policies

Cache write policies dictate the behavior of the cache when it receives a write request. There are two main write policies:

Write-through

The write is acknowledged only when the data are written both on the cache and on the slower medium.

Write-back

The write is acknowledged when data are written in cache. The slower medium is later updated with the correct data, when they need to be flushed or replaced by new ones.

Moreover, there are also policies that affect the behavior of a cache in write miss scenarios:

Write allocate

Cache loads the corresponding data block from the slower medium and the data are written to it.

No-write allocate

The write request bypasses the cache and writes directly to the slower medium.

Although the above policies can be combined as we wish, there are two main combinations that make more sense: i) write-back with write-allocate, so that writes are cached to benefit from subsequent reads and ii) write-through with no-write allocate, because the extra load of write allocate will not benefit us, since we write directly to the slower medium.

4.3.2 Caching limitations

Fast mediums like RAM and SSD drives cost more dollars/GB than slower mediums, such as hard disks **FIXME**: prove it. For this reason, caches always have smaller capacity than the mediums they cache. So, when a cache reaches its maximum capacity, it must evict one of its entries. However, which entry is the one that must be evicted?

This is a very old and well documented problem that still troubles the research community. It was first faced when creating hardware caches (the L1, L2 CPU caches we are familiar with). In 1966, Lazlo Belady proved that the best strategy is to evict the entry that is going to be used more later on in the future[Bela66]. However, the clairvoyance needed for this strategy is a little difficult to implement, so we resort to one of the following, well-known strategies:

- **Random:** Evict a randomly chosen entry. This strategy, although it seems simplistic at first, is sometimes chosen due to the ease and speed of its implementation. It is preferred in random workloads where freeing quickly space for an entry is more important than the entry that will be evicted.

- **FIFO (First-In-First-Out):** Evict the entry that was inserted first. This is also a very simplistic approach as well as easy and fast. Interestingly, although it would seem to produce better results than Random eviction, it is rarely used since it assumes that cache entries are used only once, which is not common in real-life situations.
- **LRU (Least-Recently-Used)** Evict the entry that has been less recently used. This is the most common eviction strategy and there has been a paper that uses a Bayesian statistic model to prove its optimality (**FIXME:** add paper). It is not however simple to implement since there needs to be some sort of tracking of the time of reference for each entry.
- **LFU (Least-Frequently-Used)** Evict the entry that has been less frequently used. There have been many derivatives of this algorithm that also use parts of the LRU algorithm which have promising results, but this algorithm itself is not commonly used. The reason is because it over-estimates the frequency of references to an item and it performs poorly in cases when an item is frequently accessed and then is not used at all.

The fact that write requests that have spawned the evictions cannot continue until the dirty data have been safely written to the storage backend, means that when the cache has no space left, its speed deteriorates to the speed of the slower medium.

The above observation indicates that the challenge for caching algorithms is how friendly their flushes are to the underlying storage. To elaborate on that a bit, hard disks excel on sequential payloads, so if a cache could flush in a more sequential way to the disk, it would boost its performance in these scenarios.

4.4 Real-life scenario

Usually, when a small deployment makes its first steps, it doesn't use SSDs due to management/hardware costs and since it is an investment that is actually needed when the deployment has proved that it will attract traffic. Instead, the most common setup is an array of RAID-protected commodity hard disks or fast SAS drives.

When the storage demands start to increase and more users use the service, the OS caching system of the storage nodes will soon prove ineffective and the randomness in the requested data will skyrocket the access times.

At this point, the administrators must take one (or more, if the budget allows it) of the following decisions:

1. Add more storage nodes in order to lower the load on the existing ones (horizontal scaling).
2. Buy battery-backed array controllers with volatile memory on-board, to improve access times (vertical scaling).
3. Put time-critical storage operations, such as journaling, in higher tiers (tiering)
4. Add RAM or SSD caches in write-back mode that will ACK the requests before they reach the slower mediums (caching).

The employment of one of the aforementioned techniques (scaling, tiering, caching) is of paramount importance for the future of the service.

4.5 Current Solutions

For the thesis purpose, we have evaluated a numerous of caching solutions. The results of our evaluations are presented below:

4.5.1 Bcache

Overview

Bcache has been designed by Kent Overstreet since 2011 and has been included in the Linux kernel (3.10) since the May of 2013.

Bcache allows one to use one or more fast mediums as a cache for slower ones. Typically, the slow medium is a RAID array of hard disks and the fast medium are SSD drives. Bcache has been specifically built for SSDs and has the following characteristics:

1. The data are written sequentially and in erase block size granularity, in order to avoid the costly read-erase-modify-write cycle.
2. It takes special care to mitigate wear-leveling by touching equally all SSD cells
3. It honors TRIM requests and uses them as hints for its garbage collection.

Installation and usage

Bcache is a kernel driver that needs a patched kernel and intrusive changes to the backing device.

On a nutshell, bcache edits the superblock of both the cache and backing devices in order to use them, rendering existing data unreadable. Then, it exposes to the user a virtual block device, which can be formatted to any file-system. This virtual block device is the entry point to the bcache code. Then, the caching device is attached to the backing device and at this point the virtual block device is ready to accept requests.

At any point, the bcache parameters can be further tuned via the sysfs interface.

Features and limitations

The most striking bcache feature is that it uses a custom built B+tree as an index, which has the added benefit that dirty data can be coalesced and flushed sequentially to the slower spinning medium. This provides a considerable performance speed-up for hard disks.

Some other noteworthy features of bcache are the following:

1. It can be used to cache more than one devices
2. It can operate in three modes, write-through, write-back and write-around, which can be switched on/off arbitrarily during normal usage or when the fast medium is congested.
3. It utilizes a journal log of outstanding writes so that the data are safe, even when an unclean shutdown occurs.
4. It can bypass sequential IO and send it directly to the backing device, since this workload is tailored for spinning disks.

4.5.2 Flashcache

Overview

Flashcache has been designed by Facebook and has been open-sourced in the April of 2010. It is a kernel module that is officially supported for kernels between 2.6.18 and 2.6.38 and is based on the Linux Device Mapper, which is used to map a block device onto another.

Installation and Usage

Flashcache's installation is not system-intrusive, in the sense that it needs only to compile the module against the kernel's source, modprobe it and then map the cache device upon the backing device, without making any changes to the latter.

Features and limitations

Flashcache uses a set-associative hash table for indexing. It has three modes of operation, write-through, write-back and write-around, and some basic performance tuning options such as eviction strategies and dirty data threshold. Also, it has the following limitations:

1. It does not provide atomic write operations, which can lead to page-tearing.
2. It does not support the TRIM command.

4.5.3 EnhanceIO

Overview

EnhanceIO has been developed by STEC Corp. and has been open-sourced in the December of 2012. It is a fork of Flashcache which does not use the Linux Device Mapper and has some major re-writes in parts of the code such as the write-back caching policy.

Installation and Usage

The installation method is similar to the Flashcache's method. The source code is compiled against the kernel's source, which produces a module that can be modprobed. After that, the utilities provided can be used to map the cache device on the backing device.

Features and Limitations

Similarly to Flashcache, EnhanceIO uses a set-associative hash table for indexing. It also has improvements upon the original Flashcache implementation in the following areas:

1. The page-tearing problems have been solved.
2. Dirty data flushing using background threads.

4.5.4 Memcached

Overview

Memcached is a distributed memory caching system that is being widely employed by large sites such as Youtube, Facebook, Twitter, Wikipedia. It has been created in 2003 by Brad Fitzpatrick while working in LiveJournal and to date there have been numerous forks of the code, most notably including Twitter's twemcache and fatcache, Facebook's implementation etc.

When memcached came into existence, many social sites like LiveJournal were experiencing the following problem:

User pages would often have queries that would be executed hundreds of times per second or would span across the database due to a big SELECT, but whose nature would be less critical or would not change rapidly. Queries such as "Who are my friends and who of them are online?", "What are the latest news in my feed?" etc. which could be easily cached, crippled instead the database by adding a lot of load to it.

To tackle this problem, administrator can instruct memcached to utilize the unused RAM of the site's servers to cache these kinds of queries. Ten years later, memcached has become the defacto scale-out solution, and has use cases such as Facebook's whose dedicated memcached servers serve the 95% of their queries (**FIXME:** add presentation)

Installation and usage

Memcached adheres to the client server model, with N clients connecting to M servers. Memcached, which is a user space daemon, runs on every server and listens for requests typically on port 11211. The installation is very easy since there are packages for most known distros. Once memcached has been installed, the administration needs to specify only the port and several performance options such as cache size and number of threads.

The clients on the other hand communicate with the memcached servers using native libraries. There are libraries that are written for most programming languages such as C, PHP, Python, Haskell etc. The clients can then specify which queries - or keys in general - want to be cached and the actual caching is done in runtime.

Features and limitations

Architecturally, memcached tries to do everything in O(1) time. Each memcached server consists of a hash table that indexes the keys and their data. Since the data size can vary from 1 byte to 1MB, they are organized in SLABs in order to prevent memory fragmentation. Moreover, each memcached must be able to handle tens of thousands connections from clients, so it relies in libevent to do the asynchronous polling.

What's more interesting about memcached is that its main strength is actually its biggest limitation. Memcached has no persistence and in fact, data can be evicted in numerous ways:

1. Cached data have an expiration time after which they are garbage-collected.
2. Data can be evicted before their expiration time, if the cache has become full.
3. When memcached is out of SLAB pages **FIXME:** what are slab pages?, it must evict one in order to regain space. This leads to the eviction of more than one keys.

4. When adding or removing memcached servers, the Ketama algorithm that maps keys to servers will assign a portion of the existing keys to other servers. This change in mapping, however, will not actually move the existing keys to these servers and the data are essentially invalidated.

To sum up, the lack of persistence means that memcached will never hit the disk bottleneck due to flushes and will always be very fast, as long as the cache hit rate is high. On the other hand, its unreliable nature means that it is not a general purpose software and only specific workloads will be benefited from it.

4.5.5 Couchbase Server

Overview

Couchbase server, a NoSQL database which has been under active development by Couchbase Inc. since the January of 2012, is actually the product of the merge of two independent projects, CouchDB and Membase, with CouchDB continuing as an Apache funded program. Couchbase aims to combine the scalability of memcached with the persistence of a database such as CouchDB.

Installation and usage

Couchbase provides two versions, a community edition, that lacks the latest bug fixes, and an enterprise edition. The community edition has an open-source license and can be installed easily in all major distributions from the official packages.

Once Couchbase Server has been installed, it can be configured through a dedicated web console or the command-line or the REST API. Its configuration has to do with the amount of RAM it will use and most importantly the cluster that it will join. Clusters are a deviation from the classic memcached architecture. They are logical groups of servers (or nodes) that are used for replication and failover reasons.

Like memcached, the communication with the servers is done through client libraries. These libraries are written for many different programming languages such as C, Python, Java, PHP, Ruby etc.

Features and Limitations

Couchbase Server adds the following important features to memcached feature list:

1. It can provide persistence for the data.
2. It uses data replication, which is one of the persistence guarantees.
3. It re-balances the data on resizes, so that they are evenly distributed across the database.

4.5.6 Honorary mentions

Repcached

Repcached is a memcached 1.2 fork that aims to provide asynchronous data replication. It didn't catch up however for the following reasons:

1. The added data replica merely slims the margins of losing the data but not erases them.
2. It is based in memcached 1.2, which has been released four years ago. Since then, there have been numerous performance improvements.
3. The synchronization cost of replication was high.

Ramcloud

RAMCloud (**FIXME:** add paper) is a project that is being directed by John Ousterhout at Stanford University. RAMCloud is not a caching system per se, but a storage system that uses RAM as its primary storage and hard disks for failover scenarios. To be able to stay persistent even after a node crashes, it also maintains data replicas and requires Inband connection between nodes to migrate the data quickly to other nodes.

RAMCloud states that it, or a system with similar design, will be the primary storage system of services such as AWS or Microsoft's Azure. For the time being, RAMCloud is under heavy development and its requirements are way of the budget of most deployment.

4.5.7 Evaluation

The above solutions fall into two broad categories; block store and key-value store. Both of these categories can be used in Archipelago, since there are peers that use block store semantics, e.g. when xsegbd receives a request, and peers that use key-value store semantics, e.g. when vlmc has translated a block request to object request.

We will start our evaluation with the block store techniques first. These methods have in common that they are kernel modules which cache requests that are targeted to a single (or more) slow block device. So, we can use them in this way:

1. Add SSDs to the host machine where the VMs are running.
2. Partition the SSDs so that there is one partition for each volume that is running in this host.
3. Install the kernel module and when a VM is created, run the necessary commands to map a partition of the SSD to the virtual block device of the VM.
4. Use the block device that the kernel module exposes and pass it to the hypervisor.

The main issue with this approach (and host caching in general) is this: If the cloud software is distributed then, when a host crashes, the VMs can be restarted in another host. This is possible in deployments where the instances' attributes are known by the respective cloud management software and their data are stored in a distributed storage system.

The issue arises when caching in write-back mode, where the VM's most recent data will be down with the host. In this case, the VM will not be able to start in another host or worse, it will be in inconsistent state with whatever implications this may have.

Even if we ignored the above, there are other issues too, such as:

1. If a user process segfaults, it can be restarted promptly, without interrupting the rest of the VMs. If however the kernel segfaults, the host will go down.

2. Caching at xsegbd level does not take advantage of the fact that large parts of a VM's volume are shared between other volumes due to Copy-on-Write. This means there will be lost space in the SSD for data that are actually duplicate.
3. Flashcache has page tearing issues, which we want to avoid.
4. Bcache is for newer kernels and to date it still has some bugs.
5. Having a fixed partition for each volume does not scale, since for each VM with high activity, there can be 10 other stale VMs that practically eat up cache space.

We will continue with the second category, the key-value store solutions. The programs that fall in this category have two important advantages:

1. They are distributed by nature and try to eliminate any SPOF ¹, in the same way that RADOS does.
2. They can utilize the extra RAM of a node, which is plenty in the RADOS nodes.

However, there are also some fundamental problems with them:

1. Memcached has no concept of persistence. Not only that, it basically relies on the fact that has no persistence that hacking our way through that issue would create a different software.
2. Couchbase Server has no way to use RADOS as its backing device and has its own concept of replication.

For the above reasons, we have decided to roll out our own implementation, which is presented in the following chapter.

TODO: Add why we don't want to use the page cache.

¹ Single Point of Failure

Chapter 5

Design of cached

TODO: Link the previous chapter and the evaluation of different solutions with this one

Having these observations in mind, we can provide some more strict requirements that our solution must have:

1. **Nativity:** Our solution must be native to Archipelago i.e. not need any translation layers to communicate with it.
2. **Pluggability:** Our solution must be able to provide a caching layer between peers that are already in operating mode without restarting Archipelago. Also, it must be removed without disturbing the service.
3. **In-memory:** Our solution must cache requests in RAM, since the next fastest tier, SSDs, are already being used in RADOS as a journal.
4. **O(1) complexity** Since our solution will keep data in-memory, the complexity of the indexing mechanism should be as small as possible,

For the following sections, we will drop the "*solution*" moniker and we will use instead the proper name of our implementation, "cached", which simply means **cache daemon**).

The following two chapters are the main bulk of this thesis and they present our own implementation that aims to fill the above requirements.

More specifically, this chapter provides an in-depth description of the design of cached. Section 5.1 provides the design rationale of cached and explains how its design meets the above requirements. Section 5.2 presents the building blocks of cached while Sections 5.2.2, 5.2.4 and 5.2.5 provide a detailed explanation of their design. Moreover, Section 5.2.3 illustrates the request flow of one of the most important components of cached, the xcache. Section 5.3 explains how cached utilizes the aforementioned components and introduces some unique components that have been tailored specifically for cached. Finally, in Section 5.4 we illustrate the flow of requests for cached.

5.1 Design rationale

One of the first architectural decisions was to implement cached as an Archipelago user-space peer (see Section ?? about Archipelago peers). This choice was the most natural one since it provides the smallest possible communication overhead with the other Archipelago peers. Also, this design decision covers the **nativity** requirement we posed at the beginning of this chapter.

The above design choice has another advantage too; we can plug on-line the cached peer between the vlmcd and blocker and unplug it when we want to. This opens up numerous possibilities such

as plugging cached for QoS¹ reasons, when there is a peak in I/O requests. This also means that the **pluggability** requirement is also being met.

The next important design decision was what will cached index. Given that it will reside between the `vlmcd` and `blocker`, where the VM's requests have already been translated to object requests, the natural choice is to cache objects. To understand why our caching peer must be close to the Archipelago's logic, i.e. the translation of blocks to objects, consider the following two points:

- Like `bcache`, our implementation must not only cache object requests fast but also try to coalesce them so that, when needed, they will be flushed to the slower medium in a more sequential fashion. The fact however that the VMs' volumes are partitioned into different objects, means that sequential data (in volume context) which reside in different objects will probably not be sequential in the storage backend too.

Thus, unlike `bcache` which expects that the backing device is also the physical device and coalesces data accordingly, our implementation is limited only in coalescing data in the object range (commonly 4MBs). If our implementation was caching in block or volume level, it would be unaware of that fact.

- **TODO:** Add that we are also aware of the CoW operations and can speed them up

Having decided that cached will cache objects, the next step is to decide i) on the index mechanism and ii) on **what** exactly will we index.

As for what we will index, it would be an overkill to further partition the objects and index the regions within them. Moreover, this would make sense only if the objects were large (e.g. like volumes). So, we index object names solely.

As for the index mechanism, we have chosen to use a very fast in-memory hash table that will keep the data in a preallocated space in RAM. This covers the **in-memory** requirement that we have set above. Also, the choice of the hash table is one of the reasons that our implementation has **O(1) complexity**.

Finally, another important decision was whether cached would be a multi-threaded peer. We have decided that we will implement it this way and then evaluate the performance of the implementation to find out if we are benefited by multi-threading or not.

Thus, cached must be able to work with multiple threads which will accept requests from cached's request queue and serve them concurrently with the other threads. Of course, multi-threading can be very tricky, especially when we are dealing with I/O requests and simultaneous accesses to the same object blocks. So, in order to achieve a balance between safety and speed, we use a fine-grained locking scheme in critical sections which is discussed in detail in Section 5.2.4.

5.2 Cached components

At this point, we must do an intermission before we show the design of cached. Specifically, we will show first the design of the cached's components, since many cached operations rely on them and the reader needs prior knowledge of them to grasp the cached design.

¹ Quality of Service

5.2.1 Overview

In this section, we will list the main components that cached relies on. Per Archipelago policy, most of these components have been written in the xtypes fashion (see Section ?? about xtypes).

The components of cached can be seen below:

- xcache, an xtype that provides indexing support, amongst many other things
- xworkq, an xtype that guarantees atomicity for execution of jobs on the same object
- xwaitq, an xtype that allows conditional execution of jobs

and their design will be discussed in-depth in the following sections.

Also, we must note that the above components predate our cached implementation and are not a contribution of this thesis². They are presented however in this thesis for clarity reasons.

5.2.2 The xcache xtype

xcache is the most important component of cached. It is responsible for several key aspects of caching such as:

- entry indexing,
- entry eviction,
- concurrency control and
- event hooks

Below, we can see a design graph of xcache:

FIXME: add Figure here

TODO: add better design explanation

As we can see above, xcache utilizes two hash tables. One hash table is responsible for indexing entries (or more generally speaking "cache entries") that are active in cache. The other hash table is responsible for indexing evicted cache entries that have pending jobs. Again, more generally speaking, evicted cache entries are entries whose refcount has not dropped to zero yet.

On the following subsections, we present the features of xcache as well as their design.

Entry Preallocation

Since xcache indexes a bounded number of entries, there is no need to allocate them on-the fly using malloc/free. Considering that we are caching at RAM level and not at SSD level, the system call overhead will have a considerable impact on performance. Thus, in our case, we preallocate the necessary space in advance.

² xcache is an exception since we have extended its functionalities for our purposes

Entry indexing

The index mechanism that xcache uses is a hash table named xhash, also an xtype. The reason why a hash table is used as an index mechanism is because:

1. Given that we index only a certain number of entries, we expect that the insert, lookup and delete operations are in constant time (see below for an explanation why)
2. Hash tables can preallocate the space needed whereas tries/b-trees/bst will allocate nodes as new entries are inserted. Again, the fact that we index a certain number of entries means that we expect that we will have many evictions and insertions.
3. We don't need to do substring matches (advantage of tries)
4. We don't need to traverse the entries sequentially (advantage of B-trees and BSTs)

The hash table that is used is heavily based on dictobject[dict], the Python dictionary implementation in C. Dictobject has been created to minimize the collisions and the hops (**FIXME**: Explain that better). Its only drawback is that it needs to resize when the table's entry history has reached the 2/3 of its capacity.

Besides the hash table, which answers to the question "Where is the entry?" we also need another mechanism to answer the question "Is the entry still referenced?". xcache has such a mechanism which is commonly called "reference counting". Specifically, each entry has a counter that is incremented/decremented when a user accesses/releases an entry.

To sum up, when an entry is inserted, we use its name as a key and we update its refcount to 2, one reference from the user and one standard reference from the hash table. When we lookup for an entry, we use the entry's name as a key and then increment by 1 its refcount.

Entry eviction

The decision to have xcache index a bounded number of entries means that when it reaches its maximum capacity and is requested to index a new entry, it has to resort to the eviction of a previously cached entry. Evicted entries are not removed immediately from xcache. They are instead set in an "evicted" state and they reside in a special-purpose hash table until the user confirms that they can be removed.

xcache handles evictions in an interesting way. More specifically, evictions occur implicitly and not explicitly, meaning that the user (peer) doesn't have to evict entries manually. For example, when a user tries to insert a new entry to an already full cache, the insertion will succeed and the user will not be prompted to evict an entry manually. Moreover, the user will be notified via specific event hook that is triggered upon eviction.

The scheme of implicit evictions and later on notification of the user, has the advantage that lookups, inserts and evictions can occur atomically by xcache. This wouldn't be the case if the user was responsible for the evictions.

As for the eviction strategy, we have utilized an LRU queue. Not only it's optimal (**FIXME**: verify it) for our purposes, but we have also mitigated the cost of keeping the last references for each entry by creating a simple LRU algorithm, which has $O(1)$ complexity for all update actions. More about the implementation of the LRU algorithm can be found in Section [6.1.5](#).

Concurrency control

The concept of concurrency control has been discussed in Section 2.3. The goal of xcache is to handle safely - and preferably fast - simultaneous accesses to the shared memory.

In order to do so, we must first identify which are the critical sections of xcache, to wit, the sections where a thread modifies a shared structure. These sections are the following:

- **Most xhash operations:** Inserts and removals can modify the hash table (e.g. they can resize it, add more entries or delete existing ones). This also means that lookups must not run simultaneously with the above two operations.
- **Cache node claiming:** Before an entry is inserted, it must acquire one of the pre-allocated nodes from the cache-node pool and we must ensure that this can happen concurrently from all threads.
- **Entry migration:** An entry can migrate from one hash table to the other e.g. on cache eviction. This migration involves a series of xhash operations; removal from one hash table and subsequent insertion to the other. These two operations must occur atomically.
- **Reference counting:** Every entry must have a reference counter. Reference counters provide a simple way to determine when an entry can be safely removed. Since many threads can have access to the same entry, we must provide a way to update the reference counters atomically.
- **LRU updates:** Most actions that involve cache entries must subsequently update the LRU queue. The updates at the LRU queue must also occur atomically.

Let's see what guarantees we provide for each of the above scenarios:

- **xhash operations:** We provide a lock for each hash table. Only one thread can access each hash table at any time.
- **Cache node claiming:** The cache-node queue is also protected by a lock.
- **Entry migration:** When an entry is migrated from one hash table to the other, we always acquire the lock of the hash table of active entries first and then the lock of the hash table of the evicted entries. The order on which we take the locks is very strict to avoid deadlocks.
- **Reference counting:** For the atomic increases and decreases of a counter, we don't need a lock and its added overhead. Instead, we can use the atomic get and atomic put operations that the CPU provides.
- **LRU updates:** Since the majority of LRU updates take place when a new entry is inserted in the hash table, we can protect our LRU under the same cache lock.

Event hooks

Since xcache is created to provide core caching functionalities for other peers, it must also notify them when it takes an implicit action that the peer is not aware of. In Section 5.2.2 we have seen one implicit action that xcache takes when a user inserts an entry, namely eviction.

Besides this event, there are others. The complete list is the following:

cache node initialization: This hook is triggered when a cache node is initialized. It is triggered once only for each node, during the initialization phase of xcache.

cache entry initialization: This hook is triggered when a cache entry has been inserted in the cache.

cache entry eviction: This hook is triggered when a cache entry has been evicted from the cache.

cache entry reinsertion: This hook is triggered when an evicted entry has been reinserted in the cache.

cache entry finalization: This hook is triggered when an evicted entry's refcount has dropped to 0. This serves as a warning for the user who has the opportunity to let the cache entry go or increment its refcount.

cache entry put: This hook is triggered when an evicted entry has been totally removed from the cache.

cache entry free: This hook is triggered when a removed entry's cache node has been sent back to the cache node pool.

For each of the above events, we have created the respective event hook. The peer that uses xcache may choose, if it wants, to use them and if so, it can plug its own event function for each hook which will be called when the event is triggered.

5.2.3 xcache flow

To make the way xcache works a bit more clearer, we will see the flow for three of the main xcache operations; lookup of an entry; insertion of a new entry and removal of an entry:

Insertion

FIXME: add figure and explanation

Lookup

FIXME: add figure and explanation

Put

FIXME: add figure and explanation

5.2.4 The xworkq xtype

The xworkq xtype is a useful abstraction for concurrency control. Its purpose is to enqueue "jobs" (protected by a lock) and ensure that only one thread will execute them. There is no distinction as to which thread this will be, as well as no execution condition. The executive thread is simply the one that acquires the lock first.

xworkq is generally used when multiple threads want simultaneous access to a critical section. Instead of spinning indefinitely, waiting for a thread to finish, they can enqueue their job in the xworkq and resume processing other requests. xworkq is also generic by nature, since the "job" is simply a target function and its input data.

On the following figure, we can see the design of xworkq:

FIXME: add figure

It consists of a queue where jobs are enqueued. The thread that enqueues a job can attempt to execute it too, by acquiring a lock for the xworkq. If the lock is unheld, the thread will acquire it and will be able to execute the enqueued job. Else, it can safely leave and its job will be executed by the thread that holds the lock.

In cached context, every object has an xworkq. Whenever a new request is accepted/received for an object, it is enqueued in the xworkq and we are thus ensured that only one thread at a time can have access to the object's data and metadata.

5.2.5 The xwaitq xtype

The xwaitq xtype bears some similarities to the xworkq xtype. Like xworkq, it is also an abstraction where "jobs" are enqueued and dequeued later on to be executed. Unlike xworkq though, jobs are executed only when a predefined condition is met. Another distinction is that the jobs in xwaitq are assumed to be thread-safe and can be executed concurrently by any thread.

xwaitq is commonly used in non-critical code sections that can be executed only under specific, predefined circumstances. The "jobs" that are enqueued in xwaitq are the same as the jobs of xworkq.

FIXME: add figure

Unlike xworkq, before a job is enqueued, the thread can attempt to execute it by checking the execution condition. Only if the condition is **not** met does the thread enqueue the job to the queue. Before the thread leaves, it "signals" the queue and essentially rechecks the condition to ensure that it can't be executed. It can then safely leave since its job will be executed when another thread signals the queue successfully.

In cached context, xwaitqs are used to enqueue jobs which cannot be executed immediately. Common cases are when we have run out of space, when we have run out of requests etc.

5.3 Cached Design

At this point, we have discussed in length the design of the cached components. Having the above sections in mind, we can proceed with presenting how cached has been designed.

Cached has been designed mainly as the orchestrator, a peer that utilizes several different components to handle various tasks such as indexing (xcache), concurrency (xworkq) and deferred/conditional execution (xwaitq). Cached however is not limited to the above role as these components do not cover all of the needed tasks. There are several other key tasks that cached must undertake, namely:

- *Request handling*, which is how cached handles requests from other peers and sends its own.
- *Write policy enforcing*, which is the enforcing of a cache write policy (write-through, write-back).
- *Data propagation*, which controls how data changes are propagated to the slower medium.

Moreover, cached extends its repertoire using some unique components, namely:

FIXME: Add a brief description

1. Bucket pool
2. Request wrapper (Cache-IO)
3. Book-keeping utilities

We will illustrate the design of cached from two different perspectives: the operational perspective, which can be seen in Figure 5.1, and the component perspective, which can be seen in Figure ?. Moreover, we will further explain how cached manages the above new components and tasks in the following sections.

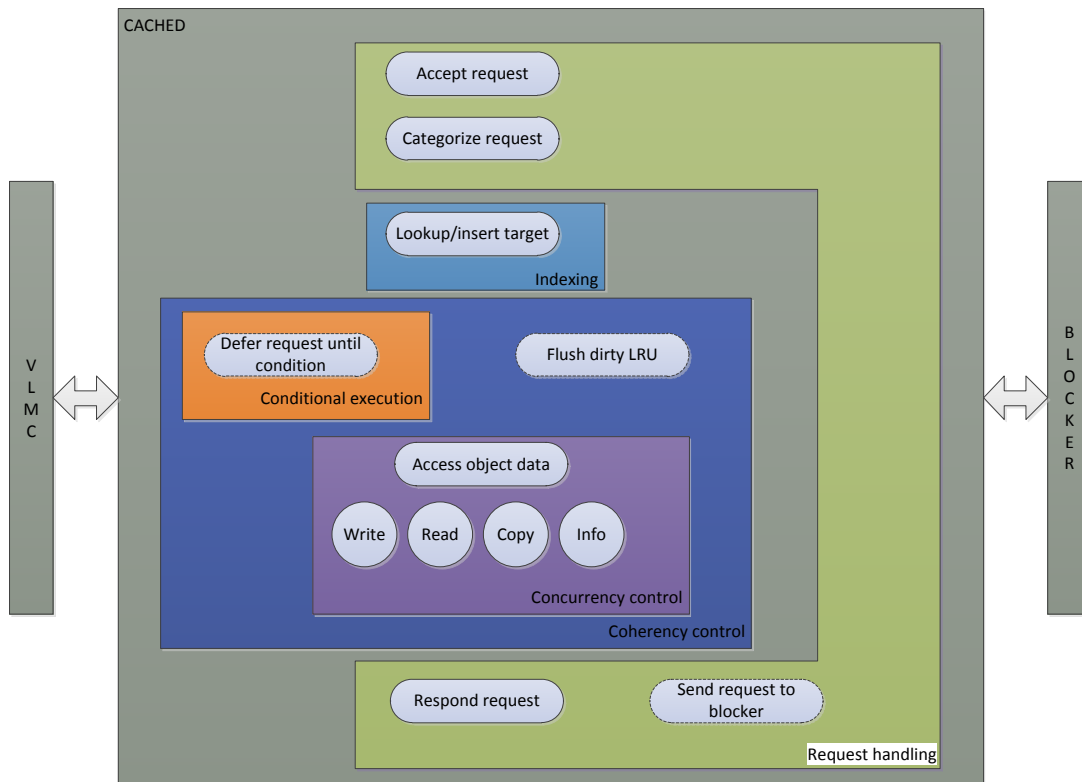


Figure 5.1: Cached design

FIXME: add component diagram for cached (xcache, xworkq, objects, buckets etc.)

5.3.1 Request handling

Cached operates as a peer that receives requests from the vlmcd. The majority of these requests will be read/write requests, but there are other types of requests too such as copy requests (sent when an object is copied-on-write) and info requests i.e. queries on what is the size of an object. Each of these requests must be handled independently, using special-purpose functions.

Furthermore, cached will also issue requests to the blocker mainly on two occasions: when it flushes a dirty object and when operating in write-through mode.

This means that, cached must be able to categorize requests and send them to the appropriate functions. Moreover, it must be able to create requests of its own, as well as handle cases such as running out of requests.

5.3.2 Write policy enforcing

The user defines beforehand what will the write policy of cached be. There are two options: write-through and write-back. These policies aren't new and have already been discussed in chapter 2, but let's see what these policies translate to in cached context.

- In **write-back** mode, cached caches writes, immediately serves the request back and marks the data as dirty. When a read arrives, it either serves the request with the dirty data (read-hit) or forwards the request to the storage peer and caches the answer (read-miss).

This policy is used when we want to improve read and write speed and can sacrifice data safety.

- In **write-through** mode, cached forwards writes to blocker, servers the request when blocker replies, caches the data and marks them as valid. When a read arrives, it either serves the request with the valid data (read-hit) or forwards the request to the storage peer and caches the answer (read-miss).

This policy is used when we want to improve read speed and want to make sure that no data will be lost.

These policies are specified once during cached's deployment and cannot be switched on/off later on.

5.3.3 Data propagation

In order to ensure that we have no data corruption, cached must provide the following two guarantees: i) consistency of cached data and ii) correct propagation of data updates to the storage backend.

As for the first guarantee, we have explained in section 5.2.4 that the consistency of the cached data is greatly secured using the `xworkq` as the guard of parallel accesses to the object's data.

As for the second guarantee, one might consider that flushing the dirty data of an object would be sufficient to correctly propagate data changes to the blocker. However, consider the scenario when a dirty object is evicted and then subsequently reinserted, overwritten and evicted again. The result would be two flush requests for the same data, sent to the blocker. Although the blocker guarantees that there will be no page-tearing, we cannot be sure about the order on which the data will be written.

To solve this problem, flush jobs must be deferred (using an `xwaitq`) if another flush for the same object is in flight.

5.3.4 Bucket pool

We have already explained in Section 5.1 the reason why cached caches objects. There is, however, one important design issue that we must address. This issue is how will cached perceive the object's data. To make it a bit clearer, given that an object typically has 4MB of size, what should cached do when a user requests e.g. a 16KB chunk of it?

If we perceived the object's data as a monolithic chunk, we would need to read and cache the whole object just to reply to the user's request. If the user then requests a chunk from another object, we would have to cache that object too and in the end, we would thrash our cache ³.

³ cache thrashing occurs when we aggressively cache data that is only used once and effectively leads to a snowball of evictions

The solution we propose is to further divide objects to the next and final logical entity, buckets (typically 4KB of size). Each bucket consists of its data and metadata and cannot be half-empty, or half-allocated. This way, we can also know which parts of the cached object are actually written, or are in the process of being read etc.

Thus, our solution to the hypothetical problem we have posed above is to request 16KB from the blocker, store the result in 4 buckets and then respond the request to the user.

The above answer, however, is not entirely complete. It implies that buckets are something readily available or attached to the object. Although each object *could* have its buckets pre-allocated, this would limit the objects that we can cache since that, even if the user requested only a small chunk, each object would statically need 4MB of space.

Ideally, we would like to be able to cache thousands of objects but i) allocate a much smaller amount of buckets and ii) strictly when the user requests to. To make things even faster, we would also like the buckets to be preallocated (like cache nodes in Section 5.2.2) to avoid the overhead of malloc/free system calls.

We have achieved the above by creating a bucket pool. The design of the bucket pool is the following:

FIXME: add diagram

The bucket pool size is static and has been set during initialization by the administrator. After the necessary space has been allocated, it is divided in buckets and all the bucket indexes are pushed in a lock-protected stack (xq). Then, each thread can pop bucket indexes from the pool and attach them to an object when needed. When that object is evicted, its attached buckets indexes are pushed back to the pool.

5.3.5 Request wrapper (Cache-IO)

When a request is accepted, all Archipelago peers commonly embed it in a peer request (you can see more about peer requests here ?). The peer request always holds the original request and optionally, several other fields that are of importance to the peer.

In our case, we keep in cached's peer requests a new structure called Cache-IO, on which we store the xcache handler of the request as well as the state of the request. Moreover, since the request may break internally in many others, we keep track of how many pending requests remain until the original one can be completed.

5.3.6 Book-keeping utilities

In order to know information such as when an object is in flushing state, when a bucket has been allocated or when a bucket is being read etc., cached must employ some sort of book-keeping of their states. The entities that require to track their states are:

1. **buckets:** Each bucket has two different states that must be tracked. The first is its allocation state:
 - i) free, meaning that the bucket is not allocated
 - ii) claimed, meaning that the bucket is allocated

The second is its data state:

- i) invalid, meaning that it currently holds no data

- ii) valid, meaning that it has data that correspond with the blockers data
 - iii) loading, meaning that a read request has been sent to the blocker to be filled with data
 - iv) writing, meaning that the bucket contents are being written/flushed to the blocker.
 - v) dirty, meaning that it currently holds newer data than what the blocker has.
2. **Objects:** The bucket states of an object provide a good indication of the object's status, yet not a complete one. The statuses we keep for the objects are:
- i) ready, meaning that it is ready to accept data
 - ii) invalidated, meaning that it is not ready to accept data
 - iii) flushing, meaning that it is currently flushing dirty data to the blocker
 - iv) failed, meaning that a request has failed for this object and we must stop using it
3. **Cache-IOs:** Cache-ios also have states. They are the following:
- i) accepted, meaning that the request has just been accepted
 - ii) reading, meaning that the request wants to read data
 - iii) writing, meaning that the request wants to write data
 - iv) served, meaning that the request has been served the data it needed
 - v) failed, meaning that the request has failed

Moreover, we keep global and per-object counters of every object's bucket states. The benefits from this approach is that we can know at any time if an object (or cached in general) has a bucket in a certain state.

5.4 Cached Flow

5.4.1 Write

This is the flow for the write path:

FIXME: add diagram and explanation

5.4.2 Read

This is the flow for the read path:

FIXME: add diagram and explanation

5.4.3 Copy

This is the flow for the read path:

FIXME: add diagram and explanation

5.4.4 Info

This is the flow for the read path:

FIXME: add diagram and explanation

Chapter 6

Implementation of cached

In the previous chapter, we have discussed in length the design of cached and its components. In this chapter, we will present how the above design has been implemented. To aid us in this task, we will use code snippets from cached, and xcache and we will comment where necessary.

More specifically, Section 6.1 presents the implementation of xcache, the main cached component. Next, section 6.2 presents the implementation of cached, showcasing the structures and functions used.

6.1 Implementation of xcache

For this section, we will attempt to provide a top-down view of the xcache implementation, starting from the functions that xcache exposes to peers and moving on to the more intrinsic details, such as the concurrency control.

6.1.1 xcache initialization

In order to use xcache, the peer must first initialize an xcache structure using `xcache_init`, which can be seen in Listing 6.1.

```
1 int xcache_init(struct xcache *cache, uint32_t xcache_size,  
2               struct xcache_ops *ops, uint32_t flags, void *priv);
```

Listing 6.1: `xcache_init` definition

`xcache_init` requests the following information from the peer:

cache: Simply, an allocated xcache struct

xcache_size: The number of objects xcache will index

ops: The trigger functions for xcache's event hooks

flags: Optional flags that tune the following two things:

1. The LRU algorithm. For the cached implementation, we use the $O(1)$ LRU, but xcache also allows to use two more LRU algorithms, a binary heap ($O(\log(N))$) or an LRU array ($O(N)$).
2. The usage of the hash table for evicted entries. Although our cached implementation relies heavily on it, this does not account for all the other peers that use xcache and by default is not used.

priv: A pointer (void *) to a structure that will be returned when an event hook is triggered. As most priv fields, it is irrelevant to the xcache struct and relevant only to the top caller. We initialize it with the peer struct.

The purpose of xcache_init is to process the above data, populate the xcache struct and create the necessary entities, such as the hash table, the cache entries etc. On Listing 6.2, we can view the xcache struct and its respective fields.

```
1 struct xcache {
2     struct xlock lock;           /* Main xcache lock */
3     uint32_t size;               /* Upper limit of entries */
4     uint32_t nr_nodes;           /* Shadow entries */
5     struct xq free_nodes;        /* Unclaimed (?) entries */
6     xhash_t *entries;            /* Hash-table for valid entries */
7     xhash_t *rm_entries;         /* Hash-table for evicted entries */
8     struct xlock rm_lock;        /* Lock for rm_entries */
9     struct xcache_entry *nodes; /* Data segment */
10    struct xcache_entry *lru;     /* O(1) lru implementation-specific */
11    struct xcache_entry *mru;     /* O(1) lru implementation-specific */
12    struct xcache_ops ops;        /* Hooks */
13    uint32_t flags;               /* Flags */
14    void *priv;                   /* Pointer to peer struct */
15 };
```

Listing 6.2: Main xcache struct

Each of the above xcache struct fields is used to implement a design feature that has already been discussed in Section 5.2.2. In the following sections, we will revisit these design features and present their implementation.

6.1.2 Cache entry preallocation

When xcache is initialized, it preallocates the necessary cache entries. The relevant fields of the xcache structure for this purpose can be seen in Listing 6.3.

```
1 struct xcache {
2     ...
3     uint32_t size;               /* Upper limit of entries */
4     ...
5     struct xq free_nodes;        /* Unclaimed (?) entries */
6     ...
7     struct xcache_entry *nodes; /* Data segment */
8     ...
9 };
```

Listing 6.3: xcache struct fields for preallocated entries

The **size** field is the number of entries. The **free_nodes** is a stack where all entry indexes are initially pushed and subsequently popped when a new entry is inserted. Finally, **nodes** is the space allocated for the cache entries and where the entry indexes point to.

Moreover, the definition of the xcache_entry struct is shown in Listing 6.4.

```
1 struct xcache_entry {
2     struct xlock lock;           /* Entry lock */
3     volatile uint32_t parallel_puts; /* Concurrency control */
4 };
```



```

4     volatile uint32_t ref;                /* Reference counter */
5     uint32_t state;                      /* Evicted or active state */
6     char name[XSEG_MAX_TARGETLEN + 1];  /* Entry name */
7     xbinheap_handler h;                  /* Index in data segment */
8     struct xcache_entry *older;          /* Less(?) recent entry in LRU
queue */
9     struct xcache_entry *younger;        /* More(?) recent entry in LRU
queue */
10    void *priv;                           /* Pointer to data contents */
11 };

```

Listing 6.4: xcache entry struct

We will comment briefly on the relevant cache entry fields for this section, which can be seen in Listing 6.5. The rest of the fields will be discussed in the following sections.

```

1 struct xcache_entry {
2     ...
3     volatile uint32_t ref;                /* Reference counter */
4     uint32_t state;                      /* Evicted or active state */
5     char name[XSEG_MAX_TARGETLEN + 1];  /* Entry name */
6     xbinheap_handler h;                  /* Index in data segment */
7     ...
8     void *priv;                           /* Pointer to data contents */
9 };

```

Listing 6.5: xcache entry fields, relevant for preallocation

The description of the fields follows:

ref The reference count of the entry, initially set to zero.

state The state of the entry. It can either be ACTIVE or EVICTED and is initially set to the first.

name The name of the entry. Since we cannot know its length beforehand, we allocate as much space as possible by our segment, typically 256 characters. During initialization, the entry name is cleared out of junk values.

h The entry's index.

priv The private contents of the cache entry. On initialization, the cache node creation hook is triggered and cached initializes the private contents of cache entry with its data (more on the Section ?)

6.1.3 Cache entry initialization

Before a peer can index a new entry, it must first allocate it from the cache entry pool and then initialize it. xcache has a special function for this purpose which can be seen in Listing 6.6

```

1 xcache_handler xcache_alloc_init(struct xcache *cache, char *name);

```

Listing 6.6: Cache entry allocation/initialization function

This function attempts to claim a cache entry from `free_nodes`. Then it initializes it with the name given by the peer. Moreover, it triggers the cache entry initialization hook which cached uses to further initialize the entry.

An added benefit of this function is that it doesn't need to acquire the cache lock, so it does not slow down the indexing functions that rely on that lock.

6.1.4 Cache entry indexing

This is the core feature of xcache. In Listing 6.7, we present the fields of xcache struct that are relevant to the indexing task.

```
1 struct xcache {
2     struct xlock lock;           /* Main xcache lock */
3     ...
4     xhash_t *entries;           /* Hash-table for valid entries */
5     xhash_t *rm_entries;        /* Hash-table for evicted entries */
6     struct xlock rm_lock;       /* Lock for rm_entries */
7     ...
8     struct xcache_entry *lru;   /* O(1) lru implementation-specific */
9     struct xcache_entry *mru;   /* O(1) lru implementation-specific */
10    ...
11 };
```

Listing 6.7: xcache struct fields for entry indexing

As we have mentioned in Section 5.2.2, we utilize two hash tables, one for the cached entries and one for the evicted entries. These hash tables can be accessed from the xcache struct and are the *entries and *rm_entries respectively.

More importantly, in Listing 6.8 we can see the functions that are related to indexing and xcache exposes to the peer:

```
1 xcache_handler xcache_lookup(struct xcache *cache, char *name);
2 xcache_handler xcache_insert(struct xcache *cache, xcache_handler h);
3 int xcache_remove(struct xcache *cache, xcache_handler h);
```

Listing 6.8: Indexing functions

All of these functions need a pointer to the xcache struct. Here's a brief description of them:

xcache_lookup: Takes the target's name as an argument and searches for it in cache.

Returns on failure: NoEntry

Returns on success: the requested handler.

Note: Looks only in entries.

xcache_insert: Takes the handler of an allocated entry as an argument and uses it to index that entry.

Returns on failure: NoEntry.

Returns on success: i) the same handler or, ii) if the same entry already exists in cache, the handler of that entry.

Note: It looks up first if the entry exist in entries or rm_entries. The later case can lead to re-insertions.

Moreover, we show in Listing 6.9 the cache entry struct fields related to indexing and comment on how they are used by each function.

```
1 struct xcache_entry {
2     ...
3     volatile uint32_t ref;       /* Reference counter */
4     uint32_t state;             /* Evicted or active state */
5     ...
6     struct xcache_entry *older; /* Less(?) recent entry in LRU
    queue */
```

```

7   struct xcache_entry *younger;          /* More(?) recent entry in LRU
queue */
8 };

```

Listing 6.9: xcache entry struct relevant indexing

The commentary on the above fields follows:

ref: The reference counter of an object is increased on lookups and inserts, since it is essentially referenced in these operations.

state: The state of an object is already ACTIVE for lookup operations (or else lookup will fail). For insertions and reinsertions, it is manually set to ACTIVE.

older/younger : These fields show the neighbors of the entry in the LRU queue. The LRU queue is sorted by reference time order, so the neighbors are essentially the entries that have been referenced right before and right after our entry.

6.1.5 Entry eviction

The relevant fields for this purpose can be seen in code listing 6.10

```

1 struct xcache {
2     ...
3     struct xq free_nodes;          /* Unclaimed (?) entries */
4     xhash_t *entries;              /* Hash-table for valid entries */
5     xhash_t *rm_entries;           /* Hash-table for evicted entries */
6     struct xlock rm_lock;          /* Lock for rm_entries */
7     ...
8     struct xcache_entry *lru;      /* 0(1) lru implementation-specific */
9     struct xcache_entry *mru;      /* 0(1) lru implementation-specific */
10    ...
11 };

```

Listing 6.10: xcache struct fields for eviction

As we have mentioned in Section ??, we resort to eviction when the cache is full and new entries can't be inserted. By xcache policy, we evict the least recently used entry. The necessary fields for the doubly-linked list that we maintain for this purpose can be seen below:

```

1 struct xcache {
2     ...
3     struct xcache_entry *lru;      /* 0(1) lru implementation-specific */
4     struct xcache_entry *mru;      /* 0(1) lru implementation-specific */
5     ...
6 };
7
8 struct xcache_entry {
9     ...
10    struct xcache_entry *older;      /* Less(?) recent entry in LRU
queue */
11    struct xcache_entry *younger;    /* More(?) recent entry in LRU
queue */
12    ...
13 };

```

Listing 6.11: Doubly-linked LRU list

The last entry of the list (oldest) is usually the LRU. When an object is referenced, it can be instantly transferred to the head of the list (MRU), since we know its position via the hash table (alternatively, we would need to search all entries, which would require $O(N)$ time).

Another feature of this LRU queue is that it doesn't require timestamps, so we can avoid the unnecessary system call.

Finally, when a cache entry is evicted from the hash table, it triggers the cache entry eviction hook.

6.1.6 Entry reinsertion

We have seen in the previous section how an active entry is evicted. However, what happens when xcache receives a request for an evicted entry which hasn't been freed yet?

In this case, the entry switches state again and is inserted back to the hash table of active entries. Also, its reference counter is incremented by 1, since the xcache can now reference this entry again.

The advantage of reinsertion is that we do not stall until the evicted entry has flushed all its data and instead, we can use it immediately.

6.1.7 Concurrency control

Concurrency control is an extremely important aspect of xcache, if we want to utilize an SMP system to its full potential. Although parts of the concurrency control have already been discussed in previous sections, in this section, we will provide an in-depth explanation of how xcache implements them.

The relevant fields for concurrency control can be seen in Listing 6.12, both for the xcache and xcache_entry structs.

```
1 struct xcache {
2     struct xlock lock;           /* Main xcache lock */
3     ...
4     struct xlock rm_lock;       /* Lock for rm_entries */
5     ...
6 };
7
8 struct xcache_entry {
9     struct xlock lock;           /* Entry lock */
10    volatile uint32_t parallel_puts; /* Concurrency control */
11    volatile uint32_t ref;        /* Reference counter */
12    ...
13 };
```

Listing 6.12: Concurrency control fields

There are three main techniques xcache uses for concurrency control. The first one is the usage of locks, which is presented in Section 6.1.7. The second one is reference counting, which is presented in Section 6.1.7. Finally, the third one is a more esoteric method that counters the ABA problem and is the tracking of parallel puts to an entry, which is presented in Section ??.

Locking

With xcache, we have tried not to use a BKL ¹ type of lock, but instead use many smaller ones.

¹ BKL stands for Big Kernel Lock and was a giant lock in kernel space that inhibited the performance of SMP systems and remained until the late stages of the 2.6 Linux kernel

Specifically, we have used:

1. a lock that protects the cache entry pool from concurrent accesses. Since the only operation this lock protects is the push and pop of cache entry indexes, we expect that there will be no contention on it.
2. the `lock` lock, as seen in the `xcache` struct, which is our main lock as it protects the hash table of active entries (`entries`) from concurrent accesses. This lock is used during lookups, inserts and evictions, so it is the lock with the most contention.
3. the `rm_lock`, which protects the hash table of evicted entries (`rm_entries`) from concurrent accesses and is used during insertions, evictions and puts.
4. a lock in every entry, which is specifically used when an entry is put.

Of major importance is also the issue of deadlocking. More specifically, during inserts or evictions, we need to have access on both hash tables. If a thread acquired the lock of one hash table and another thread acquired simultaneously the lock of the other, we would have a deadlock since both would need a lock that the other thread has.

To this end, `xcache` strictly acquires the locks in the following order: `lock` → `rm_lock` → entry lock. With this policy we are sure that there will be no deadlocks.

Reference counting

Each cache entry has a volatile `uint64_t` field which is atomically get and put. The type is volatile to inform the compiler that it might be changed outside the current execution context and therefore do not cache it in a register.

Furthermore, the atomic gets and puts are executed using the GCC builtins[`gcc-`] which are shown in Listing 6.13.

```
1 type __sync_add_and_fetch (type *ptr, type value, ...)
2 type __sync_sub_and_fetch (type *ptr, type value, ...)
```

Listing 6.13: Atomic operations of GCC

The refcount model in `xcache` should be familiar to most people:

- When an entry is inserted in cache, the cache holds a reference for it (`ref = 1`).
- Whenever a new lookup for this cache entry succeeds, the reference is increased by 1 (`ref++`)
- When the request that has issued the lookup has finished with an which entry, the reference is decreased by 1. (`ref--`)
- When a cache entry is evicted by cache, the its `ref` is decreased by 1. (`ref--`)

Moreover, some common refcount cases are:

- active entry with pending jobs (`ref > 1`)
- active entry with no pending jobs (`ref = 1`)

- evicted entry with pending jobs (ref > 0)
- evicted entry with no pending jobs (ref = 0)

Unlike most refcount cases, however, the entry is not put when its refcount drops to zero. The reason is that the entry can be reinserted at any time. In the following section, we explain how we have handled that case

Entry put

The scenario of putting the entry has proved the most tricky one and deserves its own section in the concurrency control implementation.

For this scenario, we aimed to avoid the usage of our two biggest locks: `lock` and `rm_lock`. Avoiding the first one was easy since the hash table of active entries was not used. However, the same did not hold true for the `rm_lock`.

FIXME: How much are we going to explain the ABA problem?

6.1.8 Event hooks

The hooks that xcache provides to users are stored in an `xcache_ops` struct that can be seen in Listing 6.14.

```

1 struct xcache_ops {
2     void (*on_node_init)(void *cache_data, void *data_handler);
3     int (*on_init)(void *cache_data, void *user_data);
4     int (*on_evict)(void *cache_data, void *evicted_user_data);
5     void (*on_reinsert)(void *cache_data, void *user_data);
6     int (*on_finalize)(void *cache_data, void *evicted_user_data);
7     void (*on_put)(void *cache_data, void *user_data);
8     void (*on_free)(void *cache_data, void *user_data);
9 };

```

Listing 6.14: `xcache_ops` struct

The design of these hooks has been presented on Section 5.2.2. The functions that are attached to each hook return two values: 1. the private field of xcache (the peer structure in our case) and 2. the cache entry's private data (the object in our case) for which the hook was triggered.

6.2 Implementation of cached

In this section, we will present the implementation of `cached`. Information about the design of `cached` is provided in Section 5.3. Similarly to `xcache`, we will begin with the initialization process, we will continue with the request handling of `cached` and finish with presenting the challenges we faced and the solutions we implemented.

| Switch | Info |
|--------|----------------------|
| -t | Number of threads |
| -mo | Max objects to cache |
| -ts | Total cache size |
| -os | Object size |
| -bs | Bucket size |
| -bp | Blocker port |
| -wcp | Write policy |

Table 6.1: Command line arguments of cached

6.2.1 Cached initialization

We have mentioned in the previous chapters that cached can be multi-threaded, have different write policies, maximum number of objects, cache size etc. All these variables are given from command-line and used during cached initialization.

FIXME: do not present arguments as input arguments The command-line arguments can be seen in Table 6.1

and the cached structure that is initialized is presented in Listing 6.15.

```

1 struct cached {
2     struct xcache *cache;           /* xcache struct */
3     uint64_t total_size;            /* Total cache size (bytes) */
4     uint64_t max_objects;          /* Max number of objects (plain) */
5     uint64_t max_req_size;         /* Max request size to blocker (bytes)
6     ) */
7     uint32_t object_size;          /* Max object size (bytes) */
8     uint32_t bucket_size;          /* Bucket size (bytes) */
9     uint32_t buckets_per_object;   /* Max buckets per object (
10    object_size / bucket_size) */
11    xport bportno;                 /* Blocker port */
12    int write_policy;               /* Cache write policy */
13    struct xworkq workq;            /* xworkq for deferred jobs */
14    struct xwaitq pending_waitq;    /* xwaitq for when cache entry pool
15    is empty */
16    struct xwaitq bucket_waitq;     /* xwaitq for when bucket pool is
17    empty */
18    struct xwaitq req_waitq;        /* xwaitq for when we are out of
19    requests */
20    unsigned char *bucket_data;     /* allocated space for buckets (
21    bucket pool) */
22    struct xq bucket_indexes;       /* stack of bucket indexes (bucket
23    pool) */
24    struct cached_stats stats;
25    //scheduler
26 };

```

Listing 6.15: Main cached struct

Moreover, on cached initialization we also initialize xcache as well as the general xworkqs and xwaitqs.

Some of the above cached fields are the same with the command-line arguments and are self explanatory. We will briefly comment on the less obvious fields, which will be discussed in length in their respective sections.

cache: The initialized xcache struct is stored here

max_req_size: The maximum request size that can be sent to the blocker

workq: A lockless xworkq where non-critical jobs from threads who are in a critical section are enqueued (More on Section ?)

pending_waitq: A waitq for jobs that need to allocate a cache entry to continue

bucket_waitq: A waitq for jobs that need to allocate a bucket to continue

req_waitq: A waitq for jobs that need to allocate a request to continue

bucket_data: This is a allocated space whose size is the total_size of cache. This space is later split in buckets and its indexes are pushed on a stack

bucket_indexes: The stack where bucket indexes are pushed

Furthermore, during the xcache initialization that takes place inside the cached initialization, the cached node initialization hook is triggered and cached can create its objects, which are as many as the max objects. Programmatically, cached objects are called "ce"s and their structure can be seen in Listing 6.16.

```
1 struct ce {  
2     uint32_t status; /* ce status */  
3     uint32_t *bucket_alloc_status_counters; /* counters for bucket  
4     allocation status */  
5     uint32_t *bucket_data_status_counters; /* counters for bucket data  
6     status */  
7     struct bucket *buckets; /* object buckets */  
8     struct xlock lock; /* ce lock */  
9     struct xworkq workq; /* xworkq for the entry */  
10    struct xwaitq pending_waitq; /* xwaitq for pending  
    requests on the entry */  
11    struct peer_req pr; /* Pre-allocated peer request  
    */  
12 };
```

Listing 6.16: Cached entry struct

The explanation of the above fields follows:

status: The object status, as seen in Section 5.3.6

lock: The lock for the ce's and its buckets' data

workq: The xworkq that is used for concurrency control over parallel access to the ce's and its buckets' data. It uses the aforementioned lock

pending_waitq: The xwaitq that is used when a request cannot be executed due to the ce's state. It will allow job executions only when the object is not in FLUSHING state.

We have intentionally left out the bucket related fields that will be discussed in length in Section ??.

6.2.2 Bucket pool

The initialization of the bucket pool is covered in Section 6.2.1. In this section, we will explain how this bucket pool is connected with the buckets of each ce.

When the cache node initialization hook is triggered, the ce's buckets are initialized. Essentially, this means we do (once only) the following:

1. First, we allocate an array of struct buckets. The array has `buckets_per_object` length, which is typically 1024 (4MB objects / 4KB bucket size). The struct bucket is a very simple struct and is presented in Listing 6.17.

```
1 struct bucket {  
2     unsigned char *data;  
3     uint32_t flags;  
4 };
```

Listing 6.17: Bucket implementation

2. Second, we allocate two more arrays, the `bucket_alloc_status_counters` and the `bucket_data_status_counters`, whose length is the number of allocation states (2) and data states (5) respectively.
3. Third, we initialize each bucket's allocation state to FREE and data state to INVALID. The allocation and data state are stored in the `flags` section of struct bucket, which is actually a custom bit-field with support for variable field lengths.
4. Finally, we initialize all the counters to zero, besides the allocation counter for FREE buckets, which is set to `buckets_per_object` (1024), and the data counter for INVALID buckets, which is similarly to the same number.

None of the above operations, however, interact with the bucket pool. This is because we don't initially attach the bucket indexes to the ce's buckets.

The way buckets are attached to the object is analogous to the way a function maps to its address space: a large memory chunk that has previously allocated; the chunk is internally divided to smaller chunks and when the function attempts to "touch" them, it is trapped and then the buckets are mapped to the function's address space.

Similarly, when cached accepts a request for a target, the request's range is translated to bucket range. If any of the buckets within that range are not attached to the ce, the request is trapped and the needed buckets are claimed from the bucket pool.

Finally, the bucket claiming and release procedure is the following:

1. We pop a bucket index from the `bucket_indexes` stack,
2. We translate it to the actual data pointer and store it to the `data` field of the struct bucket,
3. When the bucket is released, we translate the data pointer back to the bucket index
4. We push the bucket index back to the bucket pool.

6.2.3 Request handling

Cached uses a custom loop to poll for requests. This loop follows the same principles as the common peer's loop (see Section ?) with the addition of:

1. Checks for the state of the bucket pool. If the bucket pool has been depleted, we force flush the LRU entry to acquire its buckets.
2. Periodic signals to the cached's xworkq.

When a request is accepted/received, it is forwarded to the appropriate handle function based on its xcache operation type.

More specifically, for accepted (new) requests, we index the request target (object) and store its xcache handler on the request's cio and we proceed according to its operation type. For received requests, the request's cio holds the xcache handler for the object, so we can proceed immediately according to its operation type. The way the request is handled next is documented on Section ??.

6.2.4 "ENOSPC" scenarios

FIXME: add out of requests, out of buckets, out of cache entries

Chapter 7

Performance evaluation of cached

*"There are three kinds of lies:
lies, damned lies,
and statistics benchmarks."*
Mark Twain (modernized)

It may seem as an ironic statement, considering that we are about to provide benchmark results for cached, but it's actually a valid one. In our case, we will try not to merely smear the next pages with diagrams but first explain the benchmarking methodology behind them and then provide a concrete depiction of cached's performance under various workloads.

The skeleton of this chapter is the following: Section 7.1 explains the methodology behind our measurements. Section 7.2 provides details about the hardware on which we have conducted our benchmarks. Section 7.3 presents the results of the benchmarks that we have conducted and provides in-depth explanations about each of them. Finally, Section ? is reserved for cached on a VM (**FIXME:** take these measurements).

7.1 Benchmark methodology

The benchmarks that have been executed and whose results are presented in this chapter, will be split in two categories, both of which have their own distinct goals:

The first category is the comparison between using cached **on top** of the sosd peer (sosd has been discussed here ? **FIXME:** add section) and using solely the sosd peer.

In order to effectively compare the performance of cached and sosd, we must consider the following:

1. The comparison of the two peers should try to focus on what is the best performance that these peers can achieve for a series of tough workloads.
2. The circumstances under which both peers will be tested need not be thorough but challenging. For example, it may be interesting to test both peers against sequential requests, but i) such patterns are rarely a nuisance for production environments ii) they do not stress the peers enough to provide something conclusive iii) they are out of the scope of this section as there can be many of these kinds of tests and adding them all here will impede the document's readability.
3. Both peers must be tested under the same, reasonable workload, i.e a workload that can be encountered in production environments.
4. If the peer doesn't show a consistent behavior for a workload, it must be depicted in the results.

Having the above in mind, the next step is to choose a suitable workload. This choice though is fairly straight-forward; in production environments, the most troublesome workload is the burst of small random reads/writes and is usually the most common one that is benchmarked.

One may ponder however, how many requests can be considered as a "burst" or which block size is considered as "small". Of course, there is not only one answer to this question so, we will work with ranges. For our workload, we will use block sizes ranging from 4KB to 64KB and parallel requests ranging from 4 to 16.

The second category deals solely with the inner-workings of cached and its behavior on different operation modes or states. Its aim is not to capture the performance against a tough workload, but to explain **why** this performance is observed and how each of the options affect it. In this category, we measure how threads impact the performance of cached or what impact (if any) does our index mechanism have.

Finally, in the following sections, for brevity reasons, we will talk about comparing cached and sosd. What the reader must keep in mind however is that cached is essentially the cache layer above sosd, which means that we actually test sosd vs cached over sosd.

7.2 Specifications of test-bed

The specifications of the server on which we conducted our benchmarks is the following.

| Component | Description |
|-----------|---|
| CPU | 2 x Intel(R) Xeon(R) CPU E5645 @ 2.40GHz [e564] Each CPU has six cores with Hyper-Threading enabled, which equals to 24 threads. |
| RAM | 2 banks x 6 DIMMs PC3-10600 Peak transfer rate: 10660 MB/s |

Table 7.1: Test-bed hardware specs

| Software | Version |
|--------------|----------------------|
| OS | Debian Squeeze |
| Linux kernel | 3.2.0-0 (backported) |
| GCC | Debian 4.4.5-8 |

Table 7.2: Test-bed software specs

TODO: add configuration of RADOS cluster (journal, backing storage, network connection)

FIXME: Mention that the peers were evaluated by sending requests directly at their ports.

7.3 Performance comparison between cached and sosd

As mentioned above, for our first test, we will evaluate the read and write performance of cached and sosd for a random workload with parallel requests of small size. In order to measure accurately their performance, we will use two different metrics:

Bandwidth:

Bandwidth measures the maximum throughput that the application can sustain. This metric is usually used to indicate how much I/O (from various inputs) can an application handle within a second.

Latency:

Latency is the converse of bandwidth. It is a measurement from the viewpoint of the issuer of requests and indicates the responsiveness of the tested application. It is commonly calculated as the average reply time for a series of requests.

On the following sections, we present the benchmarks that we conducted for the first category. The first benchmark, which is shown in Section 7.3.1 attempts to depict the behavior of cached during a peak of I/O requests. The second benchmark, which is shown in Section 7.3.2, illustrates the behavior of cached under continuous load. On these sections, cached is always multi-threaded and uses 4 threads.

7.3.1 Workload smaller than cache size - Peak behavior

We begin with the bandwidth performance of our peers. The write performance can be seen in Figure 7.1 while the read performance can be seen in Figure 7.2.

Before we proceed with the interpretation of the diagram results, we will briefly comment on the diagram structure. Due to the fact that the performance of the two peers differs in at least two orders of magnitude, the results would look too flat in a conventional diagram that would scale from 0 to 11000. To amend this, we have broken the y-axis of our diagrams in two parts with different scales and starting values, in order to make the comparison easier to the eye.




Figure 7.1: Comparison of bandwidth performance for writes




Figure 7.2: Comparison of bandwidth performance for reads

The initial results look very promising. For write requests, the speedup for very small block sizes (4KB - 16KB) is approximately **100x** whereas for larger ones (32KB - 64KB) it ranges from **50x to 200x**. For read requests, the speedup for very small block sizes is approximately **50x**, whereas for larger block sizes, it ranges between **20x - 75x**.

These results not only illustrate the performance gap between RAM and HDDs but also show that our implementation manages to keep its bandwidth consistently over 1 GB/s in stress scenarios. However, we must keep in mind that they only show a part of the full picture, since it does not show how cached behaves past its cache size.

Moreover, upon closer inspection the following questions arise:

1) Where is the speedup difference between reads and writes attributed to?

We have mentioned above that the speedup for writes is between 50x and 200x whereas for reads is between 20x and 75x. Although it is a large speedup in both cases, it also shows some of the shortcomings of our implementation.

We attribute the speedup difference to three factors:

1. The performance of cached is almost the same for writes and reads. This is expected behavior as the read and write paths for cached have many common parts (**FIXME**: show write/read path) ¹.
2. Cached doesn't scale much past the 16KB block size. This is an interesting observation with an unexpected answer. It may seem implausible at first, but what happens is that we are actually hitting the bandwidth limit of the RAM modules. If we consult the table 7.1, we can see that the bandwidth limit of the RAM is about ~10.7GB/s ². This limit is approached asymptotically as

¹ To be more strict, reading seems a bit faster than writing, but that is probably attributed to CPU caches and the fact that reading from RAM is a non-destructive operation in contrast to writing.

² A more careful look shows that we surpass this limit, which is expected given the fact that we have a multi channel RAM setup.

the block size increases and the index overhead decreases (more about the index overhead later in the measurements of Section ? **FIXME:** add section).

3. On reads, sosd is benefited from the existence of caches in various levels e.g. on OSD level, on RAID controller level.

To sum up, the cached's performance remains relatively the same in both reads and writes, it is merely the sosd that is getting faster in reads due to caches.

2) Why is cached's performance increased along with block size?

We expected that sosd's performance would increase proportionally to the block size, due to the rotational nature of hard disk drives, but why does this affects cached too? It is certainly not attributed to any `memcpy()` performance tricks, since we always write in bucket granularity, which means that a 16KB write is translated to 4 x 4KB writes.

From this observation, we extrapolate that the CPU operations such as indexing, job enqueueing, accepting and responding requests, that occur proportionally more times for smaller block sizes, dominate the cached's performance.

We can see this more clearly if we consult Table 7.3.

| Bandwidth (MB/s) | IOPS | Latency (usec) | Iodepth | Block size (KB) |
|------------------|------------|----------------|---------|-----------------|
| 2652.327 | 678995.620 | 5.002 | 4 | 4 |
| 4427.891 | 566769.888 | 6.176 | 4 | 8 |
| 5900.719 | 377646.042 | 9.672 | 4 | 16 |
| 6895.167 | 220645.338 | 17.179 | 4 | 32 |
| 7090.661 | 113450.584 | 34.237 | 4 | 64 |

Table 7.3: Write performance results for 2-threaded cached

Merely using two threads instead of four lessens the performance gap significantly.

3) Why cached's performance does not improve proportionally to the parallel requests?

The reason why we see only a minor increase in the performance of cached, even though it is multi-threaded, is because our locking scheme is not fine-grained enough. We have a single lock for our request queue, a single lock for most of the hash table accesses and this inevitably causes a lot of threads to spin. This slight improvement we see is mainly due to the fact that requests are effectively being pipelined while waiting each other to release its lock.

We now proceed to the latency results. The write performance can be seen in Figure 7.3 while the read performance can be seen in Figure 7.4.




Figure 7.3: Comparison of latency performance for writes




Figure 7.4: Comparison of latency performance for reads

The latency results corroborate the observations that we made for the bandwidth results. As we can see, the latency of cached is proportional to the number of parallel requests, which reinforces the assumption that our locking scheme is not granular enough, whereas sosd has only a small variance.

7.3.2 Workload larger than cache size - Sustained behavior

We now proceed to the second part of the comparison between cached and sosd. On this part, we will once again evaluate their performance against a random workload with many parallel requests. Unlike the first part though, where the cache size was the same as the workload's, on this part the cache size will be only a fraction of it. This is after all the projected usage of cached in production environments.

For this test, there are two main parameters we must take into account: the cache size and the maximum objects. These parameters have been decoupled in our implementation and we expect different results for each combination. We have chosen to measure cache sizes that start from the 1/64th of the workload's size and reach up to the 1/8th of the workload's size. The maximum objects are chosen differently (oversubscriptions on cache size etc.) they start from a 1/1 and reach up to 8/1.

In Figure 7.5 and Figure 7.6 we can see how cached performs for the above scenario. sosd's result are of-course un-affected from the cache size and maximum cache objects, and that's will be used as indicative lines for slowness, fastness, does this word even exist?

Comparing these results with the results Figure 7.1 and 7.2, there is a vast drop in the performance. Writes specifically cannot outperform the sosd. On the other hand, reads are generally faster than sosd, about 1.5x.




Figure 7.5: Comparison of bandwidth performance for writes




Figure 7.6: Comparison of bandwidth performance for reads

Why adding more objects makes us slower?

Why reads manage to remain faster?

We will accompany the above diagrams with latency results. You can see them below:




Figure 7.7: Comparison of latency performance for writes




Figure 7.8: Comparison of latency performance for reads

7.4 Performance evaluation of cached parameters

On this part, we will see what impact do different cached parameters have on its performance. We will test the following:

1. Impact of different number of threads
2. Impact of cold cache vs. hot cache
3. Impact of writeback vs. writethrough mode

Note that the tests above are run with the following parameters:

Mode Writeback

Block size 4k

Cache size Always larger than benchmark size

The above options have been chosen to isolate cached of any other factors that may alter its performance. This way, we will be able to see more clearly the that in any other

Threads

Checking the performance impact of multi-tasking is meaningless without issuing parallel requests. Therefore, for each number of threads, we will use different IOdepth and measure its performance.

The bandwidth results can be seen in Figure 7.9 whereas the latency results can be seen in Figure 7.10.




Figure 7.9: Bandwidth performance per number of threads




Figure 7.10: Latency performance per number of threads

From these results, we derive the following conclusions:

1. Our implementation is benefited from multi-threading. We achieve a major performance improvement of up to 75% when using two threads, as well as lower performance improvement for up to four threads, as the number of parallel requests increases.

2. We don't scale well past the two threads and four parallel requests.
3. Adding more than two threads degenerates significantly the performance when the number of parallel requests is small.

Finally, these results, along with the results of the first part, clearly show the dark spot of our implementation; it needs a more fine-grained locking scheme else most of the thread's time will be spent spinning for a lock.

Cold cache vs Hot cache

This scenario will attempt to evaluate the overhead of cache misses in cached against cache hits for **write** operations. Theoretically, this should account to the overhead of adding new entries to cached and consecutively, an indication of the complexity of our index mechanism.

For this reason, we have written 128K (where **K** is *1024 and **M** is *1024²), 256K, 512K and 1M objects and have measured their latency performance. We expect that the experimental results will verify the claim that our implementation is $O(1)$.

To get the most accurate results and since we want to test just the performance of our indexing mechanism, we have also used only 1 thread and only 1 IOdepth.

On Figure 7.11 we can see the results we were talking about. The major point in these results is that we can see that write latency, either of cold or warm cache, remains practically the same as the number of objects increases.




Figure 7.11: Latency performance of cold/warm cache for variable sizes

As a side note, we observe a constant decrease in latency as the number of objects increase this is not something that should be attributed to our implementation. (explain that we have used a hash table that holds 2million objects, so it is not mapped to our process's address space. When more objects are

indexed, the hash table becomes fuller and the latency of `mmap()`s is equally distributed to the objects. Else, the hash table is more scarce but the same blocks are hit, albeit not fully written, and thus the `mmap` latency is the same but distributed to less objects.)

Chapter 8

The synapsed peer

FIXME: we probably don't care about cpu utilization but about the fact that cached will go down when the host goes down too.

On the previous sections, we have evaluated the design of cached and we have reached the conclusion that it is heavily bounded by CPU. The implications of this, however, are bigger, if we consider that all the Archipelago is running in the host machine, whose CPU's are already oversubscribed. This means that cached must compete for CPU time against the running VMs, essentially defeating the QoS purpose it serves.

TODO: Better explain that if cached could run on the under-utilized RADOS nodes, it would be accessible from any host and thus would be a major step forward towards a distributed or eventually persistent cache.

To this end, we have created a network peer called synapsed (from **synapse daemon**) as a proof-of-concept, that should be able to accept/receive any type of requests and send them through the network to another Archipelago segment. We must note that this peer has not been created with high-performance in mind but its main aim is to provide the functionality needed for our purposes. As a consequence, we haven't used tools such as ZeroMQ or libevent that would boost the performance of the implementation.

More specifically, on Section...

8.1 Design of synapsed

Given that currently Archipelago is not network-aware, peers from one segment cannot know the ports of peers of another segment. Moreover, they cannot send a request to the synapsed peer and simultaneously target a peer in another segment. So, we are faced with the problem of defining the limitations of synapsed.

We propose the following solution to this problem: Each of these two segments must have at least one synapsed peer. So, when one synapsed accepts an XSEG request, it will translate it to a network request, send it to the synapsed peer of the other segment, who will finally translate it back to an XSEG request. Moreover, each synapsed peer must be attached to a peer of its segment, which will serve as the request target when synapsed accepts request.

This means that synapsed does not actually connect two remote segments but bridge two remote peers over network.

The way synapsed handles requests can be seen in Figure ?

FIXME: add figure

The main difficulty is to make synapsed listen from its port and its request queue simultaneously. We have tackled this problem in Section [8.2.2](#).

Moreover, synapsed sends requests using the standard TCP protocol. This means that it handles the marshaling of buffers, send/receive errors as well as polling for new requests all by itself.

Finally, synapsed has been designed as a single-threaded peer.

8.2 Implementation of synapsed

TODO: Explain what are the following chapters about.

8.2.1 Synapsed initialization

Synapsed is a much simpler peer than cached and requires only the following arguments.

-hp The port where it will listen for requests.

-ra The remote address of the segment. **NOTE:** it can point to the same segment, which means that synapsed provides a generic way to bridge two peers who reside in the same host but in a different segment.

-rp The remote port of the other synapsed.

-txp The XSEG port of the peer which synapsed will attach on

Having the above arguments, synapsed can create a socket and bind its port it.

8.2.2 Request polling

Synapsed typically uses the common peer loop and Archipelago IPC methods to check its XSEG ports for new requests. However, it simultaneously needs to listen for new requests on its socket. Yet, there is no way for a process to block on its socket *and* sleep using `sigtimedwait`.

The solution seems obvious; synapsed should instead block while polling for requests in its socket and, if a new request arrives at its XSEG port, its polling will be interrupted. This is a simple solution, but there is one detail we must take into account: Sleeping like so is unsafe because peers currently block the SIGIO signal in order to wake up synchronously.

Let's elaborate on it a bit: When a peer is initialized, it blocks the SIGIO signal. Using `sigtimedwait`, it can check fast and without races if a signal has been enqueued, and then it can sleep. This is not the case with `poll()` though, since if the peer sleeps with the SIGIO signal blocked, it will not wake up. On the other hand, if we unblock the SIGIO signal, we have a new set of problems:

Consider the case where we receive a signal before we go to sleep. In order to be notified that a signal has been sent so as not to sleep, we would need to utilize a signal handler that would increment a global variable which would be checked right before we go to sleep. This approach is not only racy (what happens if SIGIO is received after this global variable is checked) but it is also very slow.

For this reason, we use the `ppoll` alternative, which accepts a signal mask as an additional argument. `Ppoll` provides the guarantee that it will atomically i. replace the old signal mask of the process with

the one provided, ii. check if there are pending signals and iii. return if there are pending signals or block until a request has been received.

Solving the issue of waiting for XSEG requests while blocking on a socket is only part of the problem. We also have to solve its counterpart, which means we need to find a quick way to iterate the XSEG ports of synapsed for new requests, while listening for new requests on its socket.

We have countered this problem by checking alternately the ports of synapsed and doing a poll with zero timeout, which returns immediately.

8.2.3 Marshalling

The XSEG request cannot be sent as is, since its data and target name are stored in two different buffers. Moreover, since these buffers can have variable length, the synapsed peer that reads from its socket needs a way to know when it has finished reading all the necessary data.

This is a common problem in computer science and its solution is to serialize (or "marshal") the object that must be transferred. The marshaling method we have chosen is the following:

1. We send at the start of every transaction a fixed-length header, whose size is known to all synapsed peers, and which has the information about the length of the rest of the buffers. The header is presented on Listing 8.1

```
1 struct synapsed_header {  
2     struct original_request orig_req;    /* Address of original  
3     requests */  
4     uint32_t op;  
5     uint32_t state;  
6     uint32_t flags;  
7     uint32_t targetlen;  
8     uint64_t datalen;  
9     uint64_t offset;  
10    uint64_t serviced;  
};
```

Listing 8.1: Synapsed header

The header is simply the necessary fields of the original request that the remote synapsed peer needs to know. The most important fields are **datalen** and **targetlen** which indicate the length of the data and target buffers respectively.

2. Once the header has been sent and the remote peer has the information that it needs, it allocates a new XSEG request from its segment and fills it with the information provided by the header.
3. Finally, the remote peer reads the rest two buffers and stores them in the XSEG request it has previously allocated.

Marshaling however is usually not so simple. There are three caveats that one must take into account before attempting to serialize manually an object. They derive from the often overlooked fact that the host and the remote machine are not architecturally the same, which can lead to:

1. **Different endianness.** This means that the two machines have completely different byte order. This is commonly solved by converting all of the data to network byte order (big endian) and then convert them to the native endianness of the machine.

2. **Different type representation.** Even if two architectures have the same endianness, it is possible that they represent the same types with different number of bytes. This is most common with 32-bit and 64-bit architectures which, for example, use 4 bytes and 8 bytes respectively to represent an int.
3. **Padding.** Due to data alignment issues, the compiler may need to pad the fields of a struct. The padding is depended on the machine's architecture but also on the compiler that is used (**FIXME:** is this true?). A common solution is to "pack" the structure, i.e. to enforce that the structure will have no padding.

For synapsed, the first two caveats do not affect us, since the machines that are used are both 64-bit, little endian machines. The third caveat may also not affect us, but just to make sure, we have packed our header structures using the gcc pragma directive:

```
#pragma pack(push, 1)
```

Listing 8.2: GCC pragma pack directive

8.2.4 Send/Receive

As we have mentioned in the previous chapter, in order to send the data from one synapsed peer to the other, we must marshal them first and then unmarshal them. This commonly requires to merge all buffers into one and send that one.

In synapsed however, we have chosen a different approach; we have employed the `readv()/writev()` functions, that allow us to do vectored I/O on these buffers.

TODO: wrap up this section

8.3 Evaluation of synapsed

TODO: add a simple diagram of synapsed's performance

Bibliography

- [1122] “RFC 1122 by Internet Engineering Task Force - October 1989”, <http://tools.ietf.org/html/rfc1122>.
- [Barr10] Thomas W. Barr, Alan L. Cox and Scott Rixner, “Translation Caching: Skip, Don’t Walk (the Page Table)”, SIGARCH Comput. Archit. Newsx, vol. 38, no. 3, pp. 48 – 59, 2010.
- [Bela66] L.A. Belady, “A study of replacement algorithms for a virtual-storage computer”, IBM Systems Journal, vol. 5, no. 2, pp. 78 – 101, 1966.
- [Caul10] Adrian M. Caulfield, Joel Coburn, Todor Mollov, Arup De, Ameen Akel, Jiahua He, Arun Jagatheesan, Rajesh K. Gupta, Allan Snively and Steven Swanson, “Understanding the Impact of Emerging Non-Volatile Memories on High-Performance, IO-Intensive Computing”, in Proceedings of the 2010 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis, SC ’10, pp. 1–11, Washington, DC, USA, 2010, IEEE Computer Society.
- [dict] “Implementation of dictobject used by Python”, <http://svn.python.org/view/python/trunk/Objects/dictobject.c?view=markup>.
- [e564] “Intel(R) Xeon(R) CPU E5645 Specifications”, http://ark.intel.com/products/48768/Intel-Xeon-Processor-E5645-12M-Cache-2_40-GHz-5_86-GTs-Intel-QPI.
- [floo] “2011 Thailand Floods”, http://en.wikipedia.org/wiki/2011_Thailand_floods.
- [gcc-] “GCC builtins for atomic memory access”, http://gcc.gnu.org/onlinedocs/gcc-4.7.2/gcc/_005f_005fsync-Builtins.html.
- [hdd-] “Would you pay \$7,260 for a 3 TB drive? Charting HDD and SSD prices over time”, <http://royal.pingdom.com/2011/12/19/would-you-pay-7260-for-a-3-tb-drive-charting-hdd-and-ssd-prices-over-time/>.
- [mt] “History of Multithreading”, <http://people.cs.clemson.edu/~mark/multithreading.html>.
- [Rose12] David S.H. Rosenthal, Daniel Rosenthal, Ethan L. Miller, Ian Adams, Mark W. Storer and Erez Zadok, “The Economics of Long-Term Digital Storage”, in The Memory of the World in the Digital Age: Digitization and Preservation, September 2012.
- [Seic06] Sugaya Seichi, “Trends in Enterprise Hard Disk Drives”, Fujitsu Sci Tech J, vol. 42, no. 1, pp. 61 – 71, 2006.
- [soun] “The disk rotation speed bottleneck”, <http://www.dbms2.com/2010/01/31/the-disk-rotation-speed-bottleneck/>.