



Εθνικό Μετσόβιο Πολυτεχνείο
Σχολή Ηλεκτρολόγων Μηχανικών
και Μηχανικών Υπολογιστών
Τομέας Τεχνολογίας Πληροφορικής
και Υπολογιστών

Thesis subject

ΔΙΠΛΩΜΑΤΙΚΗ ΕΡΓΑΣΙΑ

ΟΝΟΜΑ ΦΟΙΤΗΤΗ

Επιβλέπων : Υπεύθυνος Διπλωματικής
Τίτλος Υπευθύνου

Αθήνα, Σεπτέμβριος 9999



Εθνικό Μετσόβιο Πολυτεχνείο

Σχολή Ηλεκτρολόγων Μηχανικών
και Μηχανικών Υπολογιστών

Τομέας Τεχνολογίας Πληροφορικής
και Υπολογιστών

Thesis subject

ΔΙΠΛΩΜΑΤΙΚΗ ΕΡΓΑΣΙΑ

ΟΝΟΜΑ ΦΟΙΤΗΤΗ

Επιβλέπων : Υπεύθυνος Διπλωματικής

Τίτλος Υπευθύνου

Εγκρίθηκε από την τριμελή εξεταστική επιτροπή την 9η Σεπτεμβρίου 9999.

.....
Πρώτο μέλος επιτροπής
Τίτλος μέλους

.....
Δεύτερο μέλος επιτροπής
Τίτλος μέλους

.....
Τρίτο μέλος επιτροπής
Τίτλος μέλους

Αθήνα, Σεπτέμβριος 9999

.....
Όνομα Φοιτητή

Διπλωματούχος Ηλεκτρολόγος Μηχανικός και Μηχανικός Υπολογιστών Ε.Μ.Π.

Copyright © Όνομα Φοιτητή, 9999.

Με επιφύλαξη παντός δικαιώματος. All rights reserved.

Απαγορεύεται η αντιγραφή, αποθήκευση και διανομή της παρούσας εργασίας, εξ ολοκλήρου ή τμήματος αυτής, για εμπορικό σκοπό. Επιτρέπεται η ανατύπωση, αποθήκευση και διανομή για σκοπό μη κερδοσκοπικό, εκπαιδευτικής ή ερευνητικής φύσης, υπό την προϋπόθεση να αναφέρεται η πηγή προέλευσης και να διατηρείται το παρόν μήνυμα. Ερωτήματα που αφορούν τη χρήση της εργασίας για κερδοσκοπικό σκοπό πρέπει να απευθύνονται προς τον συγγραφέα.

Οι απόψεις και τα συμπεράσματα που περιέχονται σε αυτό το έγγραφο εκφράζουν τον συγγραφέα και δεν πρέπει να ερμηνευθεί ότι αντιπροσωπεύουν τις επίσημες θέσεις του Εθνικού Μετσόβιου Πολυτεχνείου.

Περίληψη

Περίληψη της διπλωματικής.

Λέξεις κλειδιά

Λέξη-κλειδί 1, λέξη-κλειδί 2, λέξη-κλειδί 3

Abstract

Abstract of diploma thesis.

Key words

Key-word 1, Key-word 2, Key-word 3

Ευχαριστίες

Ευχαριστίες.

Όνομα Φοιτητή,
Αθήνα, 9η Σεπτεμβρίου 9999

Η εργασία αυτή είναι επίσης διαθέσιμη ως Τεχνική Αναφορά CSD-SW-TR-*-* , Εθνικό Μετσόβιο Πολυτεχνείο, Σχολή Ηλεκτρολόγων Μηχανικών και Μηχανικών Υπολογιστών, Τομέας Τεχνολογίας Πληροφορικής και Υπολογιστών, Εργαστήριο Τεχνολογίας Λογισμικού, Σεπτέμβριος 9999.

URL: <http://www.softlab.ntua.gr/techrep/>
FTP: <ftp://ftp.softlab.ntua.gr/pub/techrep/>

Contents

Περίληψη	5
Abstract	7
Ευχαριστίες	9
Contents	11
List of Figures	13
1. Introduction	15
1.1 Introduction/Motivation	15
1.2 Thesis structure	15
2. Chapter 2	17
2.1 Section 1	17
2.1.1 Subsection 1	17
2.2 Section 2	17
2.2.1 Subsection 1	17
2.2.2 Subsection 2	17
3. Chapter 3	19
3.1 Necessary theoretical background	19
3.1.1 Multi-threaded programming	19
3.1.2 IPC	19
3.2 Archipelago	19
3.3 XSEG	20
3.3.1 Drivers	20
3.3.2 Libraries	20
3.3.3 Xtypes	20
3.3.4 Peers	20
4. Tiering	21
4.1 Theoretical Background	21
4.1.1 Caching	21
4.2 Existing storage tiers	22
4.2.1 Beache	22
4.2.2 Memcached	22
4.2.3 Blabla	22
4.2.4 Summary	22
5. Design of cached	23
5.1 Design overview	23

5.1.1	And we cache... what exactly?...	24
5.1.2	Cached components	24
5.2	The xcache xtype	24
5.2.1	Entry Preallocation	25
5.2.2	Entry indexing	25
5.2.3	Entry eviction	25
5.2.4	Concurrency control	26
5.2.5	Re-insertion	27
5.2.6	xcache flow	27
5.3	The xworkq xtype	27
5.4	The xwaitq xtype	28
5.5	Cached internals	28
5.5.1	Object states	28
5.5.2	Per-object peer requests	28
5.5.3	Write policy	28
5.6	Cached Operation	29
5.6.1	Write-through mode	29
5.6.2	Write-back mode	29
6.	Implementation of cached	31
6.1	Implementation of xcache	31
6.1.1	Entry Preallocation	31
6.1.2	Entry Indexing	32
6.1.3	Concurrency control	32
Bibliography		35

List of Figures

Chapter 1

Introduction

1.1 Introduction/Motivation

Bla-bla...

1.2 Thesis structure

Chapter 2: We define what "cloud" means and mention some of the most notable examples. Then, we give a brief overview of the synnefo implementation, its key characteristics and why it can have a place in the current cloud world.

Chapter 3: We present the architecture of Archipelago and provide the necessary theoretical background (mmap, IPC) the reader needs to understand its basic concepts. Then, we thoroughly explain how Archipelago handles I/O requests. Finally, we mention what are the current storage mechanisms for Archipelago and evaluate their performance.

Chapter 4: We explain why tiering is important and what is the state of tiered storage at the moment (bcache, flashcache, memcached, ramcloud, couchbase). Then, we provide the related theoretical background for cached (hash-tables, LRUs). Finally, we defend why we chose to roll out our own implementation.

Chapter 5: We explain the design of cached, the building blocks that is consisted of (xcache, xworkq, xwaitq). Then, we give some examples that illustrate the operation under different scenarios

Chapter 6: We present the cached implementation, the structures that have been created and the functions that have been used.

Chapter ??: We explain how cached was evaluated and present benchmark results.

Chapter ??: It connects brain parts. And its tale must be told.

Chapter ??: We draw some concluding remarks and propose some future work.

Chapter 2

Chapter 2

2.1 Section 1

This section has an important citation[[aeal99](#)]

2.1.1 Subsection 1

This subsection has code in Haskell:

```
1 foo [] = []  
2 foo h:t = 9: foo t
```

Listing 2.1: Sample code

It also has a list:

Item 1 First item

Item 2 Second item and a footnote¹.

Item 3 Third item and text in *italics*.

And an enumerated list:

1. First item.
2. Second item and text in **bold**

2.2 Section 2

2.2.1 Subsection 1

This subsection has a link to the block of code [2.1](#) in Section 1.

2.2.2 Subsection 2

This subsection has a FIXME comment, visible only to the author.

¹ Footnote description.

Chapter 3

Chapter 3

3.1 Necessary theoretical background

3.1.1 Multi-threaded programming

Multi-threading programming is good and is bad and here are some challenges:

1. Concurrency control
2. Challenge 2
3. Challenge 3

Concurrency control

Locking Three concepts for locking:

1. Lock overhead
2. Lock contention
3. Deadlocking

3.1.2 IPC

Below we can see some IPC methods:

1. mmap()
2. Semaphores
3. Sockets

3.2 Archipelago

Archipelago consists of the following:

1. XSEG 2. 3.

3.3 XSEG

XSEG is the segment on which the IPC...

There are some XSEG stuff such as:

1. Drivers 2. Libraries 3. Xtypes 4. Peers

3.3.1 Drivers

3.3.2 Libraries

3.3.3 Xtypes

The rationale behind xtypes is:

- Abstraction(?) layers: Creating inner abstractions layers for software is not a new concept but it's very easy to miss, especially when you start small and end up big.

In a nutshell, when writing code for a new software (in our case a peer for Archipelago but this can apply to most software that surpass the 1000 LOC¹ mark) it is wrong practice to create from scratch a monolithic implementation with indistinguishable parts. There is a main reason for this:

Monolithic implementations usually derive from lack of code architecture and planning. Although it is feasible for a programmer to create fully-functional code that meets the necessary requirements, albeit with a lot more effort and concentration, this approach will backfire when the programmer needs to add new features. Since there is no explicit code architecture and the fragile inner correlations are between lines of code and not separate entities, stored precariously in the developer's mind, the result will eventually be constant code refactorization.

One might think that new features happen once in a while in the development cycle but that would be wrong. This happens more often than you might think and is actually the common case in iteration and test-driven development.

The right practice instead is to...

- Re-usability:...
- User-space / Kernel-space agnosticity: (I doubt that such a word even exists...)

3.3.4 Peers

¹ Lines Of Code

Chapter 4

Tiering

4.1 Theoretical Background

4.1.1 Caching

In caching, there are usually the following two policies:

- Write-through: This policy bla bla bla
- Write-back: This policy blu blu blu

Eviction

Caching generally means that you project a large address space of a slow medium to the smaller address space of a faster medium. That means that not everything can be cached as there is no 1:1 mapping. So, when a cache reaches its maximum capacity, it must evict one of its entries

And the big question now arises: which entry?

This is a very old and well documented problem that still troubles the research community. It was first faced when creating hardware caches (the L1, L2 CPU caches we are familiar with). In 1966, Lazlo Belady proved that the best strategy is to evict the entry that is going to be used more later on in the future[Bela66]. However, the clairvoyance needed for this strategy was a little difficult to implement, so we had to resort to one of the following, well-known strategies:

- **Random:** Simply, a randomly chosen entry is evicted. This strategy, although it seems simplistic at first, is sometimes chosen due to the ease and speed of each. It is preferred in random workloads where getting fast free space for an entry is more important than the entry that will be evicted.
- **FIFO (First-In-First-Out):** The entry that was first inserted will also be the first to evict. This is also a very simplistic approach as well as easy and fast. Interestingly, although it would seem to produce better results than Random eviction, it is rarely used though, since it assumes that cache entries are used only once, which is not common in real-life situations.
- **LRU (Least-Recently-Used)**
- **LFU (Least-Frequently-Used)**

Choosing the LRU strategy is usually a no-brainer. Not only does it *seem* more optimal than the other algorithms, but it has also been proven, using a Bayesian statistic model, that no other algorithm that tracks the last K references to an entry can be more optimal.

4.2 Existing storage tiers

4.2.1 Bcache

4.2.2 Memcached

4.2.3 Blabla

4.2.4 Summary

Chapter 5

Design of cached

In the previous chapters, we have addressed the need for tiering in terms of scalability as well as performance.

We have also evaluated current caching solutions and described why they couldn't be used as a cache tier in Archipelago.

With the results of chapter ? in mind, we can provide some more strict requirements that our solution must have:

1. Requirement 1
2. Requirement 2
3. Requirement 3
4. Requirement 4

The following two chapters are the main bulk of this thesis and they present our own implementation that aims to fill the above requirements.

More specifically, this chapter provides an in-depth description of the design of cached. Section ? provides a general overview of cached. Sections ? - ? present the building blocks of cached and their design. Section ? presents the interaction of cached and its building blocks. Finally, in Section ? we illustrate the flow of requests for cached.

5.1 Design overview

Cached is a peer that operates between the vlmc and blocker. Every request that the vlmc peer sends to the blocker, the cached peer can *intercept* it, so to speak, and cache it. This follows the same principle with bcache, which plugs its own `request_fn()` function to the virtual device it creates. Unlike bcache however, cached can be plugged on and off at any time.

Cached has two different caching policies, **write-through** and **write-back**. These policies aren't new and have been discussed extensively in chapter ?, but let's see what these policies translate to in cached context.

- In **write-back** mode, cached caches writes, immediately serves the request back and marks the data as dirty. When a read arrives, it either serves the request with the dirty data (read-hit) or forwards the request to the storage peer and caches the answer (read-miss).

This policy is used when we want to improve read and write speed and can sacrifice data safety.

- In **write-through** mode, cached forwards writes to blocker, servers the request when blocker replies, caches the data and marks them as valid. When a read arrives, it either serves the request with the valid data (read-hit) or forwards the request to the storage peer and caches the answer (read-miss).

This policy is used when we want to improve read speed and want to make sure that no data will be lost.

5.1.1 And we cache... what exactly?...

Since Archipelago divides internally the volumes to objects (usually 4MB), the cached peer must operate on object level. Also, in order to know which parts of the cached object are actually written, or are in the process of being read etc. cached further divides objects to the next and final logical entity, buckets (typically 4KB). Each bucket consists of its data and metadata and cannot be half-empty, or half-allocated. You can say that the bucket is the quantum of cached objects.

I'll attempt to make the above a bit clearer. When cached receives a request, it first checks the request target (i.e. the object name and then calculates which bucket objects are within the request's range. It is easy to see that this is a 1:1 mapping to the object's data.

The fact that objects are pre-allocated means two things:

1. We don't need to care about memory fragmentation and system call overhead
2. We cannot index single buckets. <FILLME>

5.1.2 Cached components

Let's see now the design of cached in detail. The cached peer consists of a number of building blocks. Per Archipelago policy, most of these building blocks have been written in the xtypes fashion. The reasons behind this decision have been discussed in chapter ? but for completeness' sake, we will mention once more the merits of xtypes:

The components of cached can be seen below:

- xcache, an xtype that provides indexing support, amongst many other things
- xworkq, an xtype that guarantees atomicity for execution of jobs on the same object
- xwaitq, an xtype that allows conditional execution of jobs
- bucket pool, a pre-allocated memory pool for buckets

and their design will be discussed in-depth in the following sections.

5.2 The xcache xtype

xcache is the main component of cached. It is responsible for several key aspects of caching such as:

- entry indexing,

- entry eviction, and
- concurrency control

Below we can see a design overview of xcache:

As we can see above, xcache utilizes two hash tables. One hash table is responsible for indexing entriess (or more generally speaking "cache entries") that are active in cache. The other hash table is responsible for indexing evicted cache entries that have pending jobs. Again, more generally speaking, evicted cache entries are entries whose refcount has not dropped to zero yet.

5.2.1 Entry Preallocation

Since xcache has a bounded number of entries that will allocate, there is no need to allocate them on-the fly using malloc/free. Considering that we are caching at RAM level and not at SSD level, the system call overhead will have a considerable impact on performance.

Thus, the best thing to do in our case would be to pre-allocate the necessary space.

5.2.2 Entry indexing

In order to index the cached entries, xcache relies on another xtype, xhash, which is a hash table. Moreover, it's actually the C implementation of the dictionary used in Python.

We have chosen to use a hash table as our index because:

Finally, the xhash xtype gives provides us with the basic hash table functions, namely:

- Insertion
- Look-up
- Deletion

5.2.3 Entry eviction

As we can see in figure ?, xcache has been designed to index a pre-defined number of entries. That means that when xcache reaches its maximum capacity and is requested to index a new entry, it has to resort to the eviction of a previously cached entry. We have chosen the LRU strategy

Also, an added bonus is that we won't need to sacrifice speed over optimality, since that, our hash table approach allows us to create an $O(1)$ LRU algorithm which you can see in the following figure:

In a nutshell, our LRU implementation uses a doubly linked list blablabla. This design allows us to do all of the following action in constant time:

- Insert a new entry to the LRU list
- Evict the LRU entry
- Update an entry's access time (i.e. mark it as MRU)
- Remove an arbitrary entry

Another interesting feature of xcache is that evictions occur implicitly and not explicitly. The user doesn't need to interact with the LRU queue.

For example, when a user tries to insert a new entry to an already full cache, the insertion will succeed and the user will not be prompted to evict an entry manually. Also, the user will be notified via specific event hook that is triggered upon eviction that an entry has been evicted.

More about hooks can be seen in the following subsection.

5.2.4 Concurrency control

The concept of concurrency control has been discussed in chapter ?. The goal of xcache is to handle safely - and preferably fast - simultaneous accesses to shared memory.

In order to do so, we must first identify which are the critical sections of xcache, that is the sections where a thread can modify a shared structure. These sections are the following

- All xhash operations: Two of the three xhash operations (inserts and removals) can modify the hash table (e.g. they can resize it and reallocate space for it). This means that the third one (lookups) must not run concurrently with the other.
- Cache node claiming: Before an entry is inserted, it must acquire one of the pre-allocated nodes and we must ensure that this can happen from all threads.
- Entry migration: An entry can migrate from one hash table to the other e.g. on cache eviction. This migration involves a series of xhash operations; removal from one hash table and subsequent insertion to the other. This a scenario that must be handled properly.
- Reference counting: Every entry must have a reference counter. Reference counters provide a simple way to determine when an entry can be safely removed. You can see more about reference counting in chapter ?
- LRU updates: Most actions that involve cache entries must subsequently update the LRU queue. Being a doubly linked list, if two threads update the LRU simultaneously, we can lead to seg-faults.

Let's see what guarantees we provide for each of the above scenarios:

- xhash operations: We provide a lock for each hash table
- Cache node claiming: The free node queue is protected by a fast lock
- Entry migration: We always take fist the lock for entries and then for rm_entries
- Reference counting: Another important guarantee is the reference counting of entrys. xcache uses atomic gets and puts to update the reference count of an entry.
- LRU updates: Since all LRU operations take place for entries in "entries" hash table and LRU updates are blazing fast we can secure our LRU with the cache->lock.

5.2.5 Re-insertion

We have previously mentioned that in xcache, there can be data migration between hash tables. This is easy to see why in case of evictions: an entry that previously was in "entries" must now be migrated to "rm_entries" until its reference count falls to zero and can be freed.

However, what happens when xcache receives a request for an evicted entry?

there is a concept called "re-insertion". In order for an entry to be re-inserted to the primary hash table (which will be called "entries" from now on) it must first reside in the hash table that indexes the evicted cache entries (which will be called "rm_entries" from now on). As mentioned above, an entry that is in rm_entries has probably pending jobs that delay its removal.

So, what happens if a lookup arrives for that entry while on this stage? In this case, we re-insert it to entries and increase its refcount by 2, since there is one reference by the hash table and one reference by the one who requested the lookup.

5.2.6 xcache flow

Below we will see three important scenarios

Insertion

Figure

Lookup

Figure

Put

Figure

5.3 The xworkq xtype

The xworkq xtype is a useful (what?) for concurrency control on object level. It is important to distinguish between cache level operations and object level operations. Cache level operations include insertions, lookups, removals, allocations and refcount handling. On object level, there is a different set of operations that must be synchronized across threads. Namely, we have bucket claiming, read/write operations and object flushes.

The above distinction makes it easy to see that provided that operations on object level need not worry about interactions with other objects. Each object is "sandboxed", so to speak.

Let's see the design of the xworkq xtype. It consist of a queue where jobs (e.g. read from block, write to block) are enqueued. The thread that enqueues a job can attempt to execute it to, by acquiring a lock for the workq. If the lock is free, the thread will be able to execute the enqueued job. Also, other threads can enqueue their jobs, so the thread that has the lock can do those too. There is an xworkq for every object.

Every object has a workq. Whenever a new request is accepted/received for an object, it is enqueued in the workq and we are sure that only one thread at a time can have access to the objects data and metadata.

For more information, see the xworkq.

5.4 The xwaitq xtype

When a thread tries to insert an object in cache but fails, due to the fact that cache is full, the request is enqueued in the xcache waitq, which is signaled every time an object is freed.

For more information, see the xwaitq.

5.5 Cached internals

5.5.1 Object states

Every object has a state, which is set atomically by threads. The state list is the following:

- **READY**: the object is ready to be used
- **FLUSHING**: the object is flushing its dirty buckets
- **DELETING**: there is a delete request that has been sent to the blocker for this object
- **INVALIDATED**: the object has been deleted
- **FAILED**: something went very wrong with this object

Also, object buckets have their own states too:

- **INVALID**: the same as empty
- **LOADING**: there is a pending read to blocker for this bucket
- **VALID**: the bucket is clean and can be read
- **DIRTY**: the bucket can be read but its contents have not been written to the underlying storage
- **WRITING**: there is a pending write to blocker for this bucket

Finally, for every object there are bucket state counters, which are increased/ decreased when a bucket state is changed. These counters give us an $O(1)$ glimpse to the bucket states of an object.

5.5.2 Per-object peer requests

Reads and writes to objects are practically read/write request from other peers, for which a peer request has been allocated. There are cases though when an object has to allocate its own peer request e.g. due to a flushing of its dirty buckets. Since this must be fast, there are pre-allocated requests hard-coded in the struct of each object which can be used in such cases.

5.5.3 Write policy

The user must define beforehand what is the write policy of cache. There are two options: write-through and write-back. On a side note, as far as reads and cache misses are concerned, cached operates under a write-allocate policy.

5.6 Cached Operation

5.6.1 Write-through mode

Here we will see how cached operates in write-through mode.

Write

This is the flow for the write path:

Read

This is the flow for the read path:

5.6.2 Write-back mode

Here we will see how cached operates in write-back mode.

Write

This is the flow for the write path:

Read

This is the flow for the read path:

Chapter 6

Implementation of cached

In the previous chapter, we presented a design overview for cached and its components. In this chapter we will blabla how the above design has been implemented and explain in depth the structures and functions that have been created for this purpose.

More specifically, sections ? - ? provide implementation information for the components of cached, as described in Chapter ?. Next, section ? presents the actual initialization and blabla operations using excerpts from the code.

6.1 Implementation of xcache

In this section we describe how we implemented the design concept of Section 5.2.

This is the main xcache struct:

```
1 struct xcache {
2     struct xlock lock;
3     uint32_t size;
4     uint32_t nr_nodes;
5     struct xq free_nodes;
6     xhash_t *entries;
7     xhash_t *rm_entries;
8     struct xlock rm_lock;
9     struct xcache_entry *nodes;
10    uint64_t time;
11    uint64_t *times;
12    struct xbinheap binheap;
13    struct xcache_ops ops;
14    uint32_t flags;
15    void *priv;
16 };
```

Listing 6.1: Main xcache struct

Each of the above xcache struct fields serves a design purpose. Let's see which fields help in what:

6.1.1 Entry Preallocation

The relevant code for this purpose can be seen in Listings ??.

```
1 struct xcache {
2     ...
3     uint32_t nr_nodes;
4     struct xq free_nodes;
```

```

5     ...
6     struct xcache_entry *nodes;
7     ...
8 };

```

Listing 6.2: xcache struct fields for preallocated entries

When entries are preallocated, they take up a contiguous space in memory. The start of this space is the where the *nodes field points to

6.1.2 Entry Indexing

The relevant code for this purpose can be seen in Listings ??.

```

1 struct xcache {
2     ...
3     uint32_t size;
4     ...
5     xhash_t *entries;
6     xhash_t *rm_entries;
7     ...
8 };

```

Listing 6.3: xcache struct fields for entry indexing

6.1.3 Concurrency control

Reference counting The refcount model in xcache should be familiar to most people:

- When an entry is inserted in cache, the cache holds a reference for it (ref = 1).
- Whenever a new lookup for this cache entry succeeds, the reference is increased by 1 (ref++)
- When the request that has issued the lookup has finished with an entry, the reference is decreased by 1. (ref-)
- When a cache entry is evicted by cache, the its ref is decreased by 1. (ref-)

Some common refcount cases are:

- active entry with pending jobs (ref > 1)
- active entry with no pending jobs (ref = 1)
- evicted entry with pending jobs (ref > 0)
- evicted entry with no pending jobs (ref = 0)

and, as always, the entry is freed only when its ref = 0.

Case	Refcount
active entry with pending jobs	$\text{ref} > 1$
active entry with no pending jobs	$\text{ref} = 1$
evicted entry with pending jobs	$\text{ref} > 0$
evicted entry with no pending jobs	$\text{ref} = 0$

Table 6.1: Reference counting of xcache

Bibliography

- [aeal99] Some author et al., “Name of citation”, in *Proceedings of the 99th ACM Symposium on Something (POPL'99)*, pp. 999–999, Nine, 9999.
- [Bela66] L.A. Belady, “A study of replacement algorithms for a virtual-storage computer”, *IBM Systems Journal*, vol. 5, no. 2, pp. 78 – 101, 1966.